



CSE 325-Operating System

Lecture: 1-3

Computer Science and Engineering





Course Description

- This course introduces the principles and techniques for the **design and implementation of operating systems**:
 - **interrupts**,
 - **computer resource management**
 - memory management,
 - processor management,
 - I/O management,
 - file management,
 - process management and
 - security management,
 - **inter process communication.**
- Additionally topics are: Multithreaded OS, and Concurrent computations.
- This course includes a project implementation.



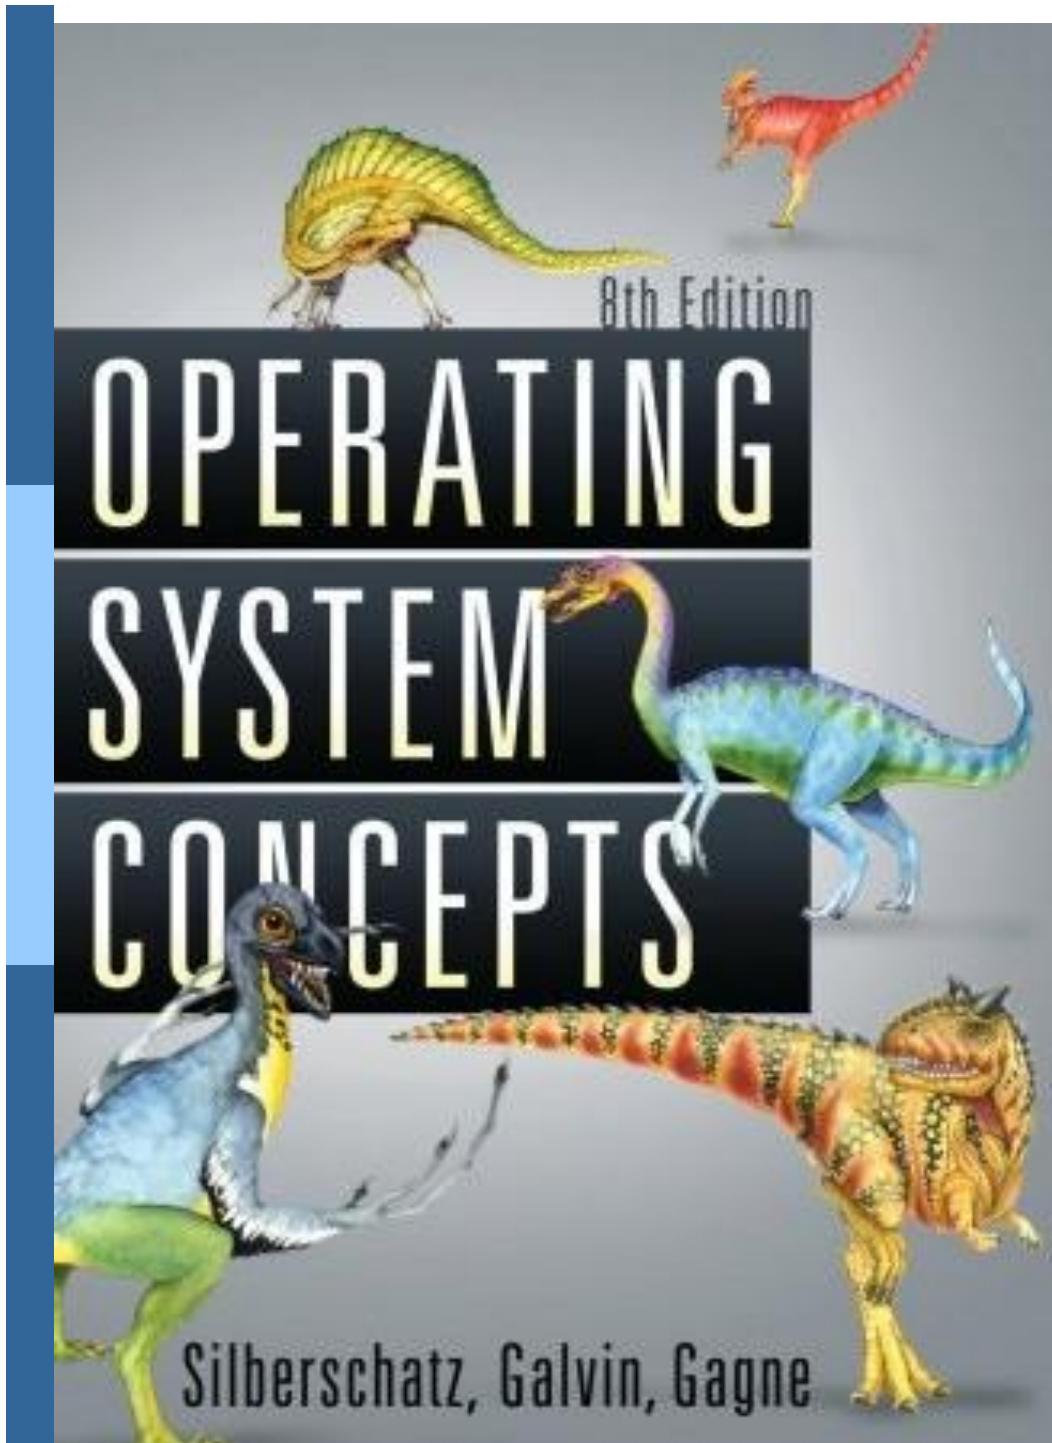


Course Outcomes

Successful completion of this course, you will be able to:

- Identify different components of Operating System.
- Describe different methods of resource management.
- Apply resource management techniques for resource constrained problems.
- Expose memory organization and I/O management techniques.





Text book



Silberschatz, Galvin and Gagne ©2009



Lecture 1: Introduction





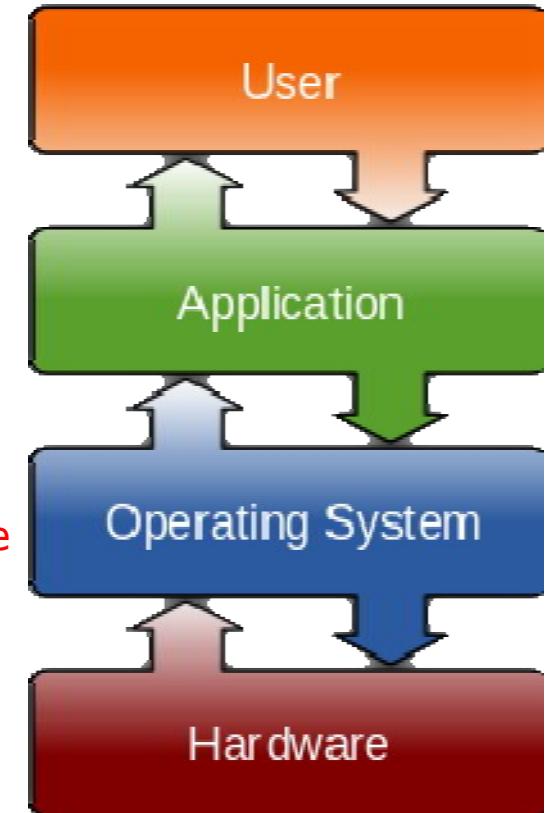
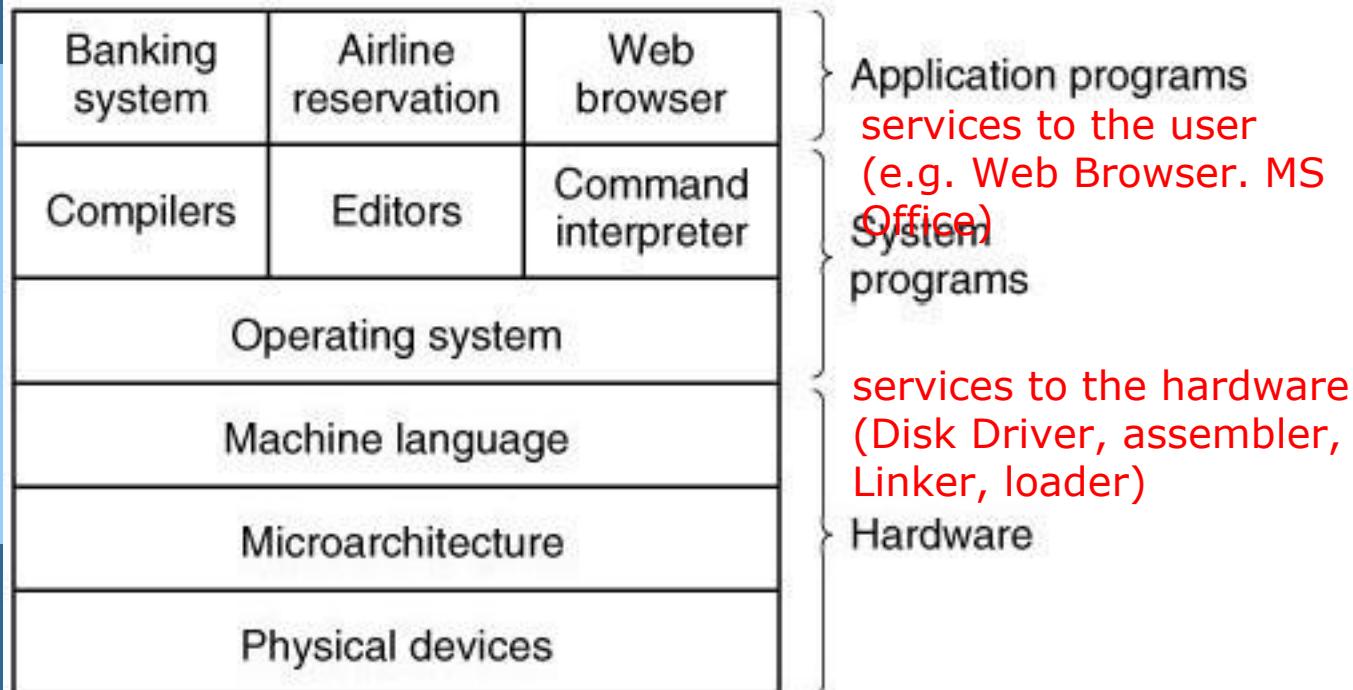
Objective

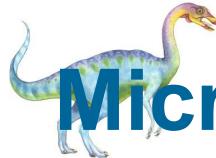
Overview of the major OS components.





Components of a Computer System





Microarchitecture- Computer Organization

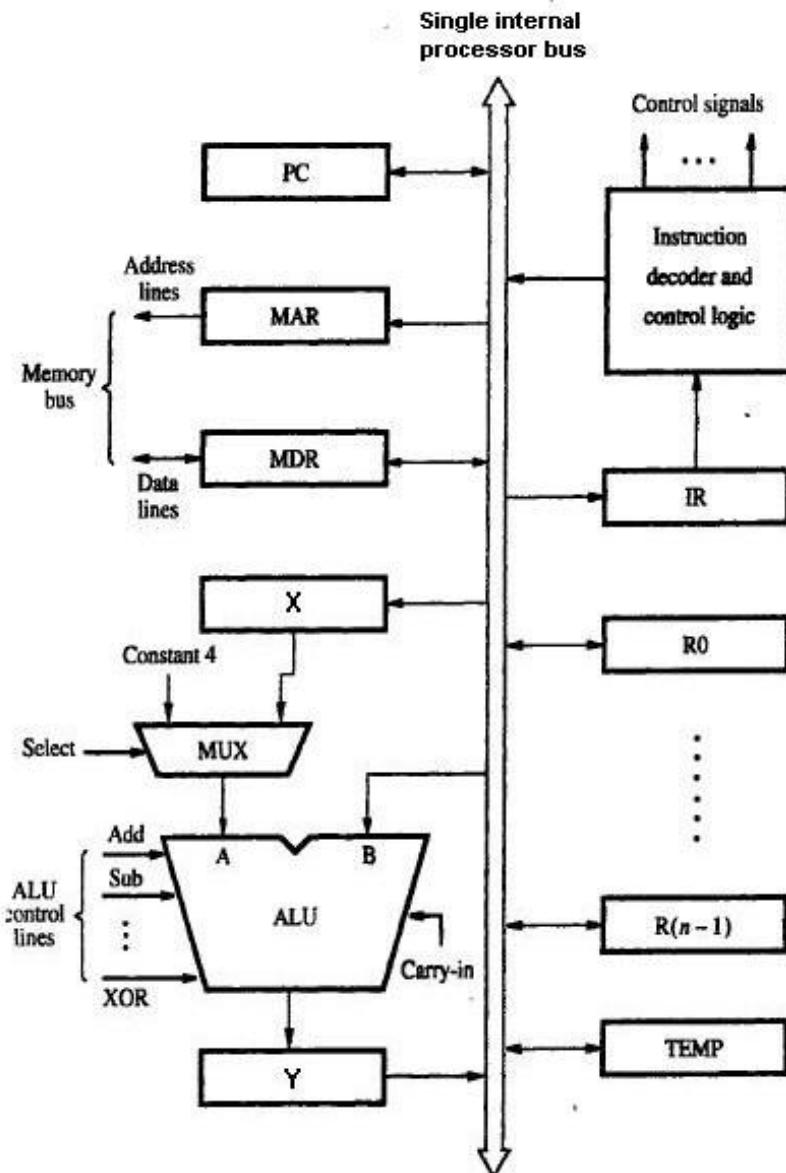
- The way a given instruction set architecture (ISA) is implemented on a processor

The ISA includes :

- the execution model,
- processor registers,
- address and data formats.

The microarchitecture includes:

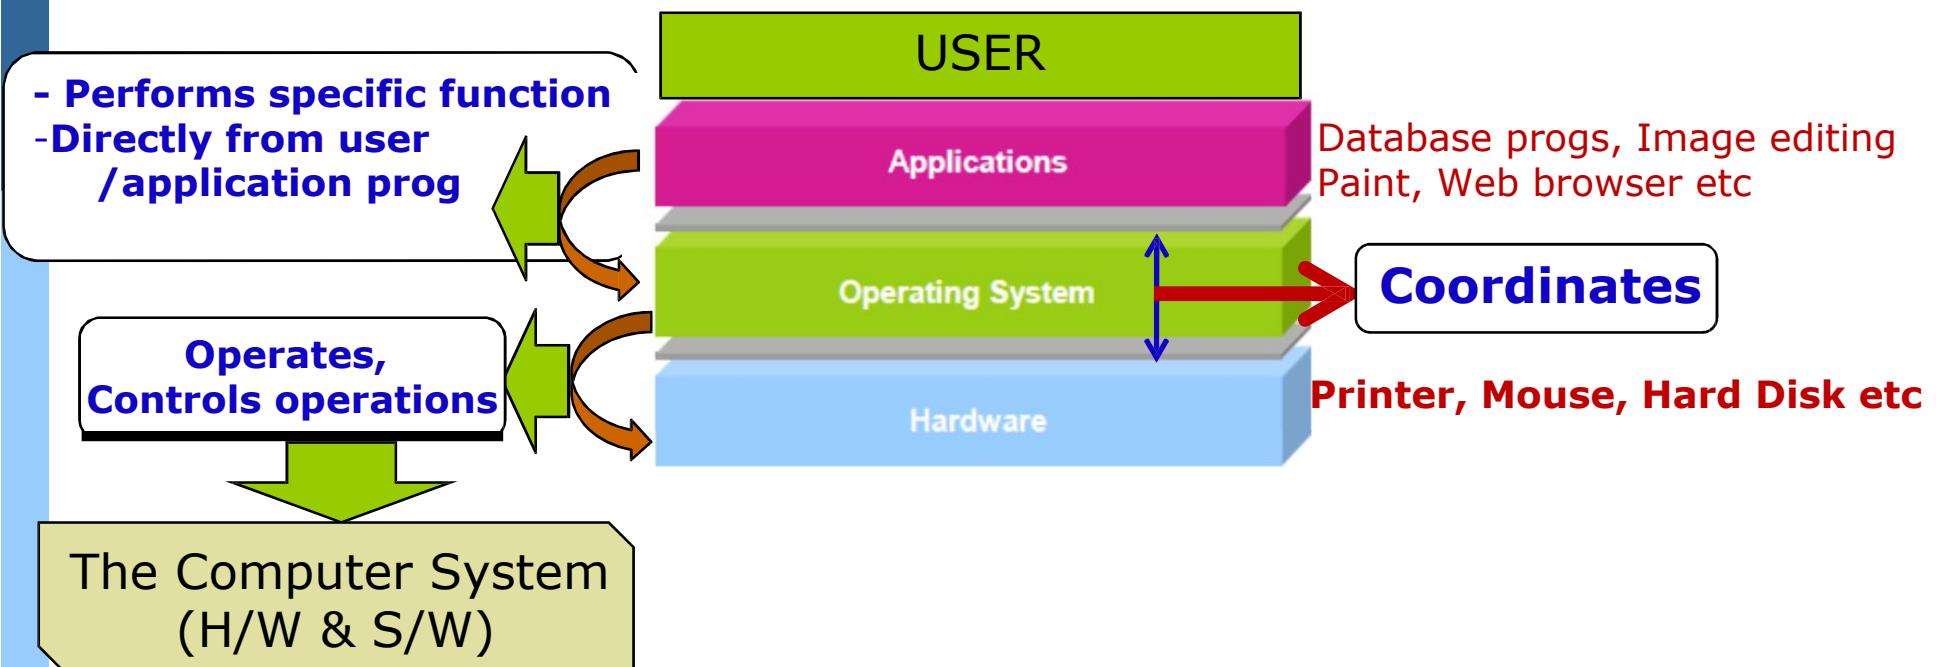
- the parts of the processor and
- how these interconnect and
- interoperate to implement ISA.





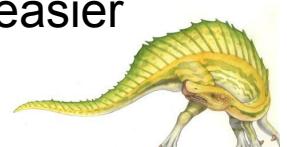
What is an Operating System?

- A program that acts as an intermediary betⁿ a user & H/W



GOALS:

- » Convenient to use
- » Execute user programs & make solving user problems easier
- » Use H/W in an efficient manner





Operating System Definition

■ OS is a **resource allocator**

- Manages all resources
- Decides between conflicting requests for efficient and fair resource use.



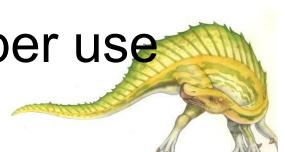
One goal of the operating system is to increase the utilization of resources.

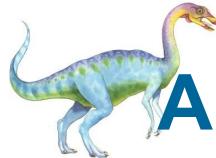
Utilization =useful time/total time

For example, the OS should avoid wasting CPU time because the disk is rotating or wasting switching among tasks. Waiting for I/O.

■ OS is a **control program**

- Controls execution of programs to prevent errors and improper use of the computer. e.g. Memory Protection.





An example comparing life with/without OS

Life with an OS

```
file = open  
    ("test.txt",  
     O_WRONLY);  
write (file, "test",  
4); close (file);
```

Life *without* an OS

- Blocks, platter, track, and sector
- Where is this file on disk? Which platter, track, and sectors?
- Code needs to change on a different system





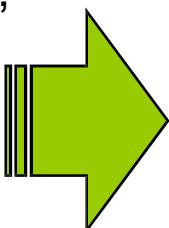
How does a program execute?

■ C Program

```
int A[2]; int C[2]; int B[2];
P=A[0];
Q=B[0];
C[0]=P+Q;
```

First.c

Compile



■ Assembly

```
Load $t0, [A]
Load $t1, [B]
ADD $t3, $t0, $t1
STORE $t3, [C]
```

First.asm

First.o

Assemble

■ Executable

0001.....00	00
0000.....01	01
00011.....11	11
100.....11	11

PC=0

IR= 0001.....00

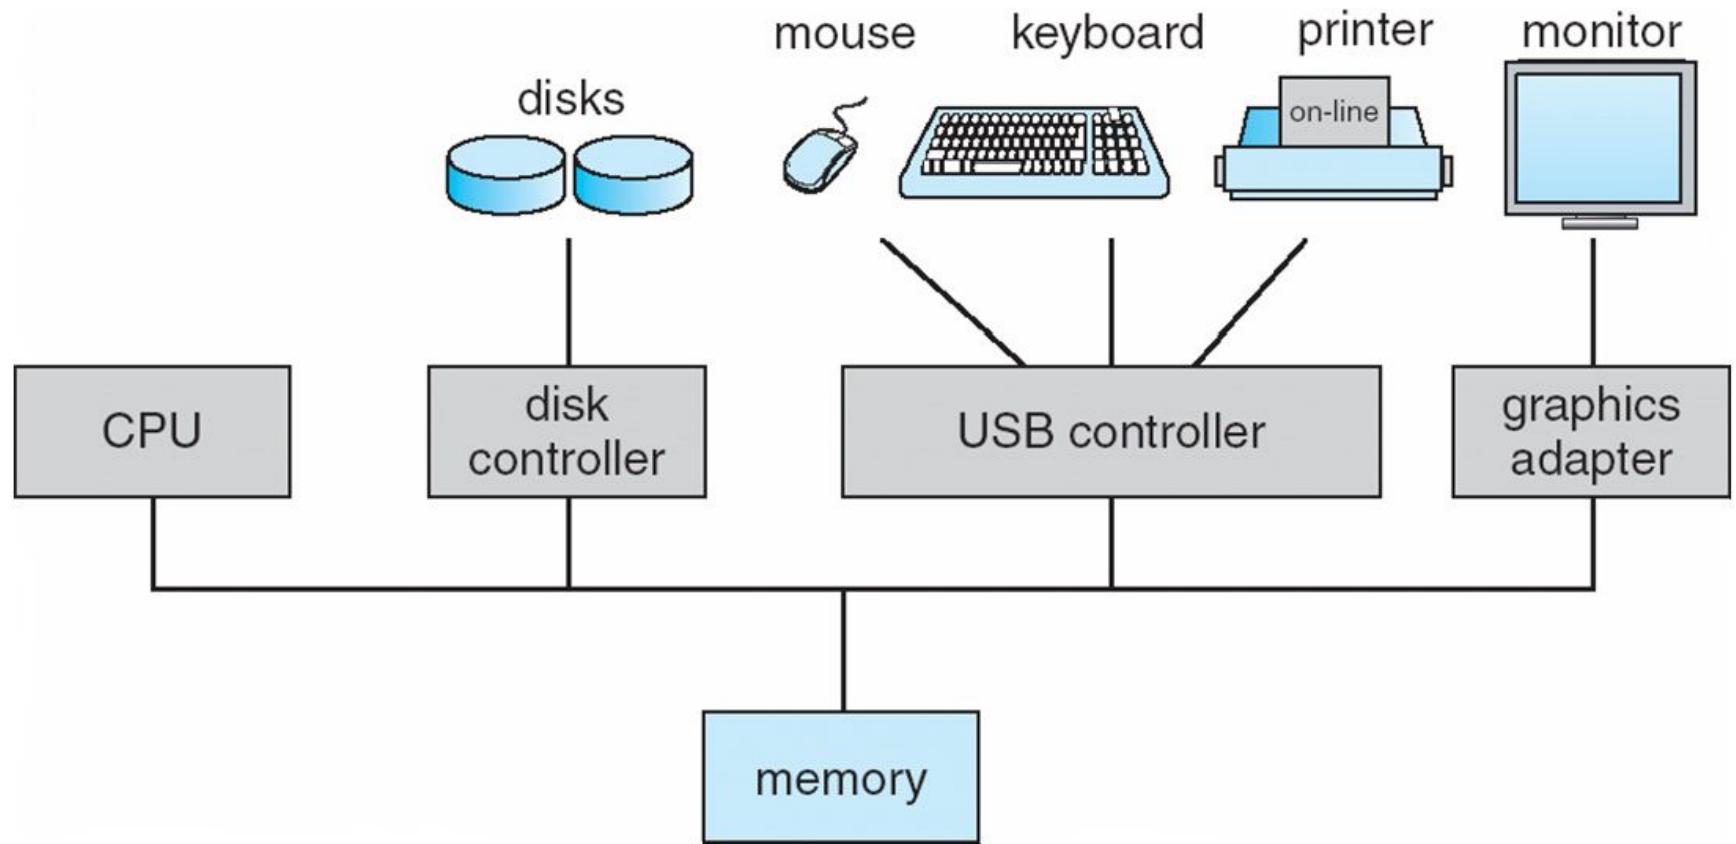
■ Memory

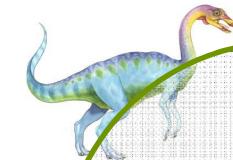
0	0001.....00	00
4	0000.....01	01
8	00011.....11	11
12	100.....11	11
100.....12	11	





Computer System





I/O

Mapped: Special I/O instruction in/out instruction

Example from the Intel architecture: `out 0x21, AL`

`000-00F DMA Controller`

`020-021 Interrupt Controller`

`040-043 Timer`

`3D0-3DF Graphics Controller`

Memory mapped I/O: No Special Instruction load/store instructions

Controller registers/memory maps in physical address space. I/O accomplished with load and store instructions

Example: `load 0xFFFF0000, AL`

(Control) `load 0xFFFF0004, BX`

1-Ready
0-not

32 bits

IE

R

32 bits Data



Polling VS Interrupt

Both are methods to notify processor that I/O device needs attention

Polling

simple, but slow

processor check status of I/O device regularly to see if it needs attention

similar to checking a telephone without bells!

Interrupt

fast, but more complicated

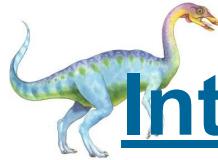
Interrupt is better if processor has other work to do

processor is notified by I/O device (interrupted) when device needs attention

similar to a telephone with bells

In brief: Interrupt is asynchronous, and polling is synchronous





Interrupts Work

Service provider

CPU

Interrupts

Status bits

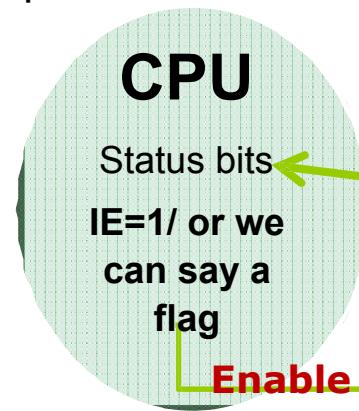
IE=0 / or we
can say a flag

**Not
enable**

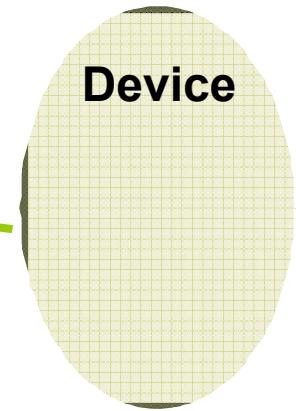
Masked Interrupts
Short periods of time

requester

Device



Interrupts



**Sends acknowledgement, disable
interrupts IE=0, and calls Interrupt Service
routine**



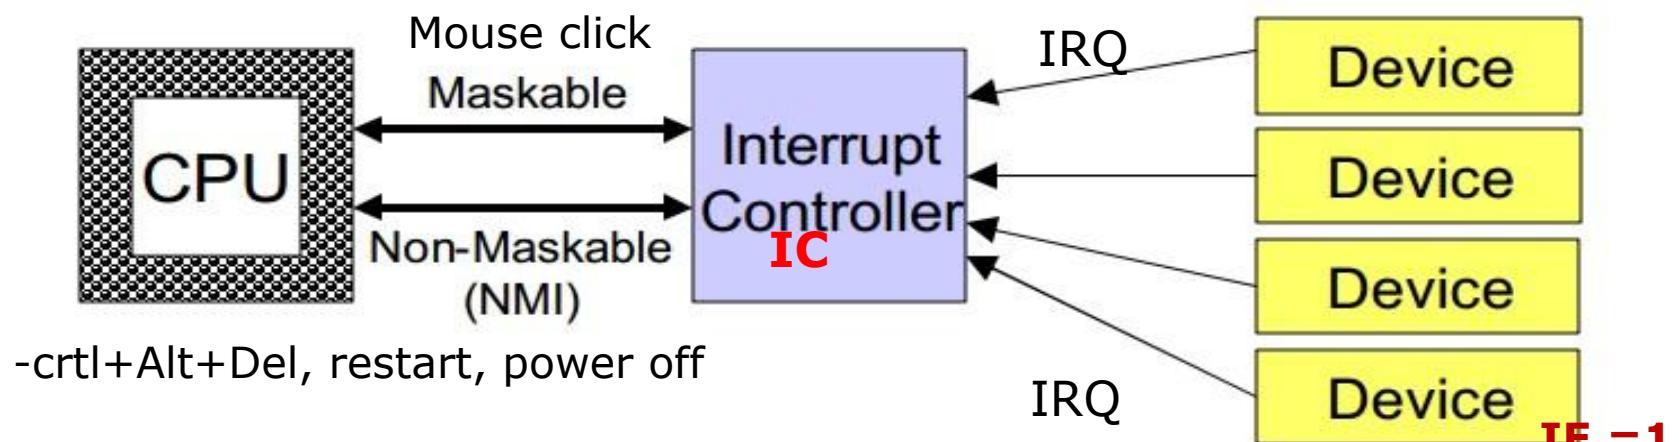
Silberschatz, Galvin and Gagne ©2009



Alternative: Interrupt Vector

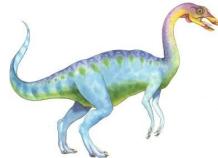
- Give each device a wire (interrupt line) that it can use to signal the processor
 - When interrupt signaled, processor executes a routine called an interrupt handler to deal with the interrupt
 - No overhead when no requests pending

IE = 1



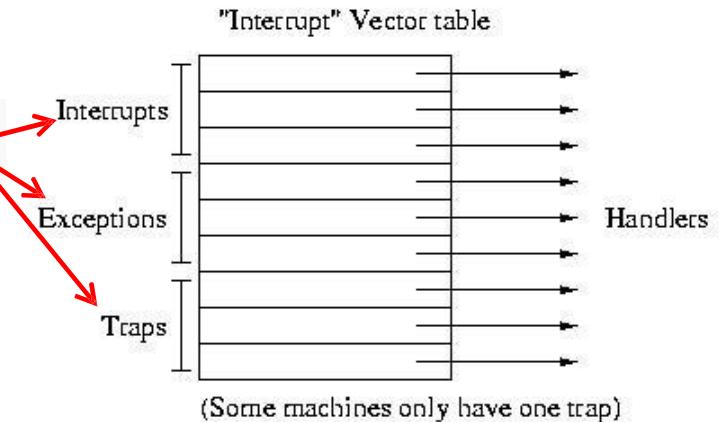
CPU sets CPU IE = 1, Device can send interrupt (Not Masked)





Different Interrupts

An operating system is interrupt driven

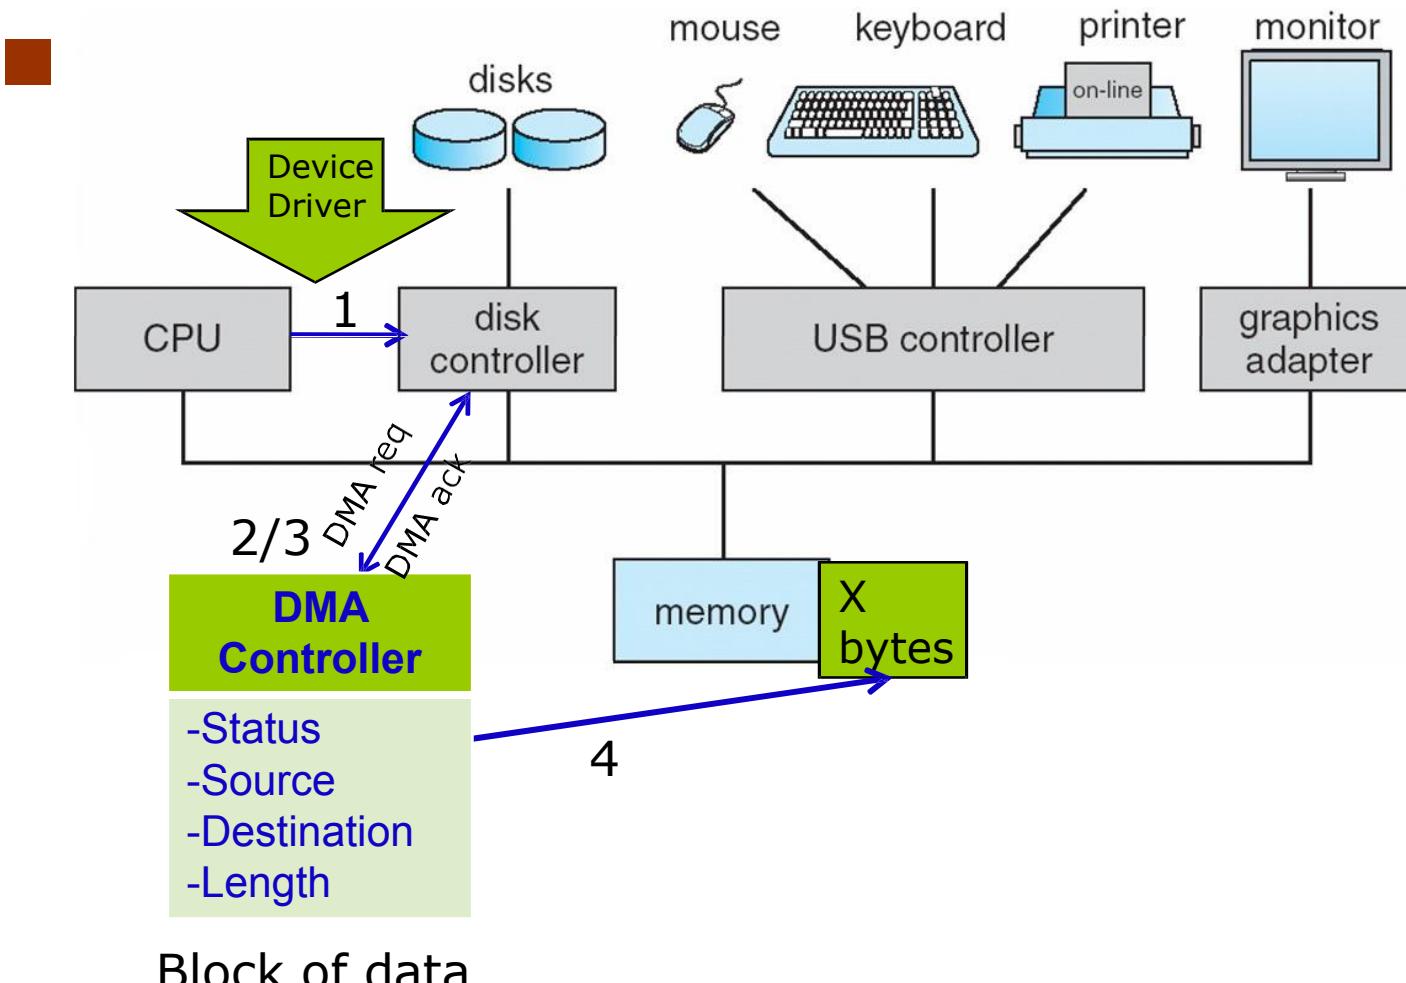


Non-Maskable (hardware)	<ul style="list-style-type: none">It does not check IE flag or priority.
Maskable (hardware)	<ul style="list-style-type: none">It checks IE flag or priority. If IE=1, the processor acknowledges the interrupt..
Normal interrupt (Software Interrupt)	<ul style="list-style-type: none">a programmer initiated and expected transfer of control to a special handler routine.user program can ask for an operating system serviceunconditional –control always transfer to pre-defined procedure.
Exception (Software Interrupt)	<ul style="list-style-type: none">automatically generated unexpected events occur in response to some exceptional condition.illegal program actions that generate an interrupt.



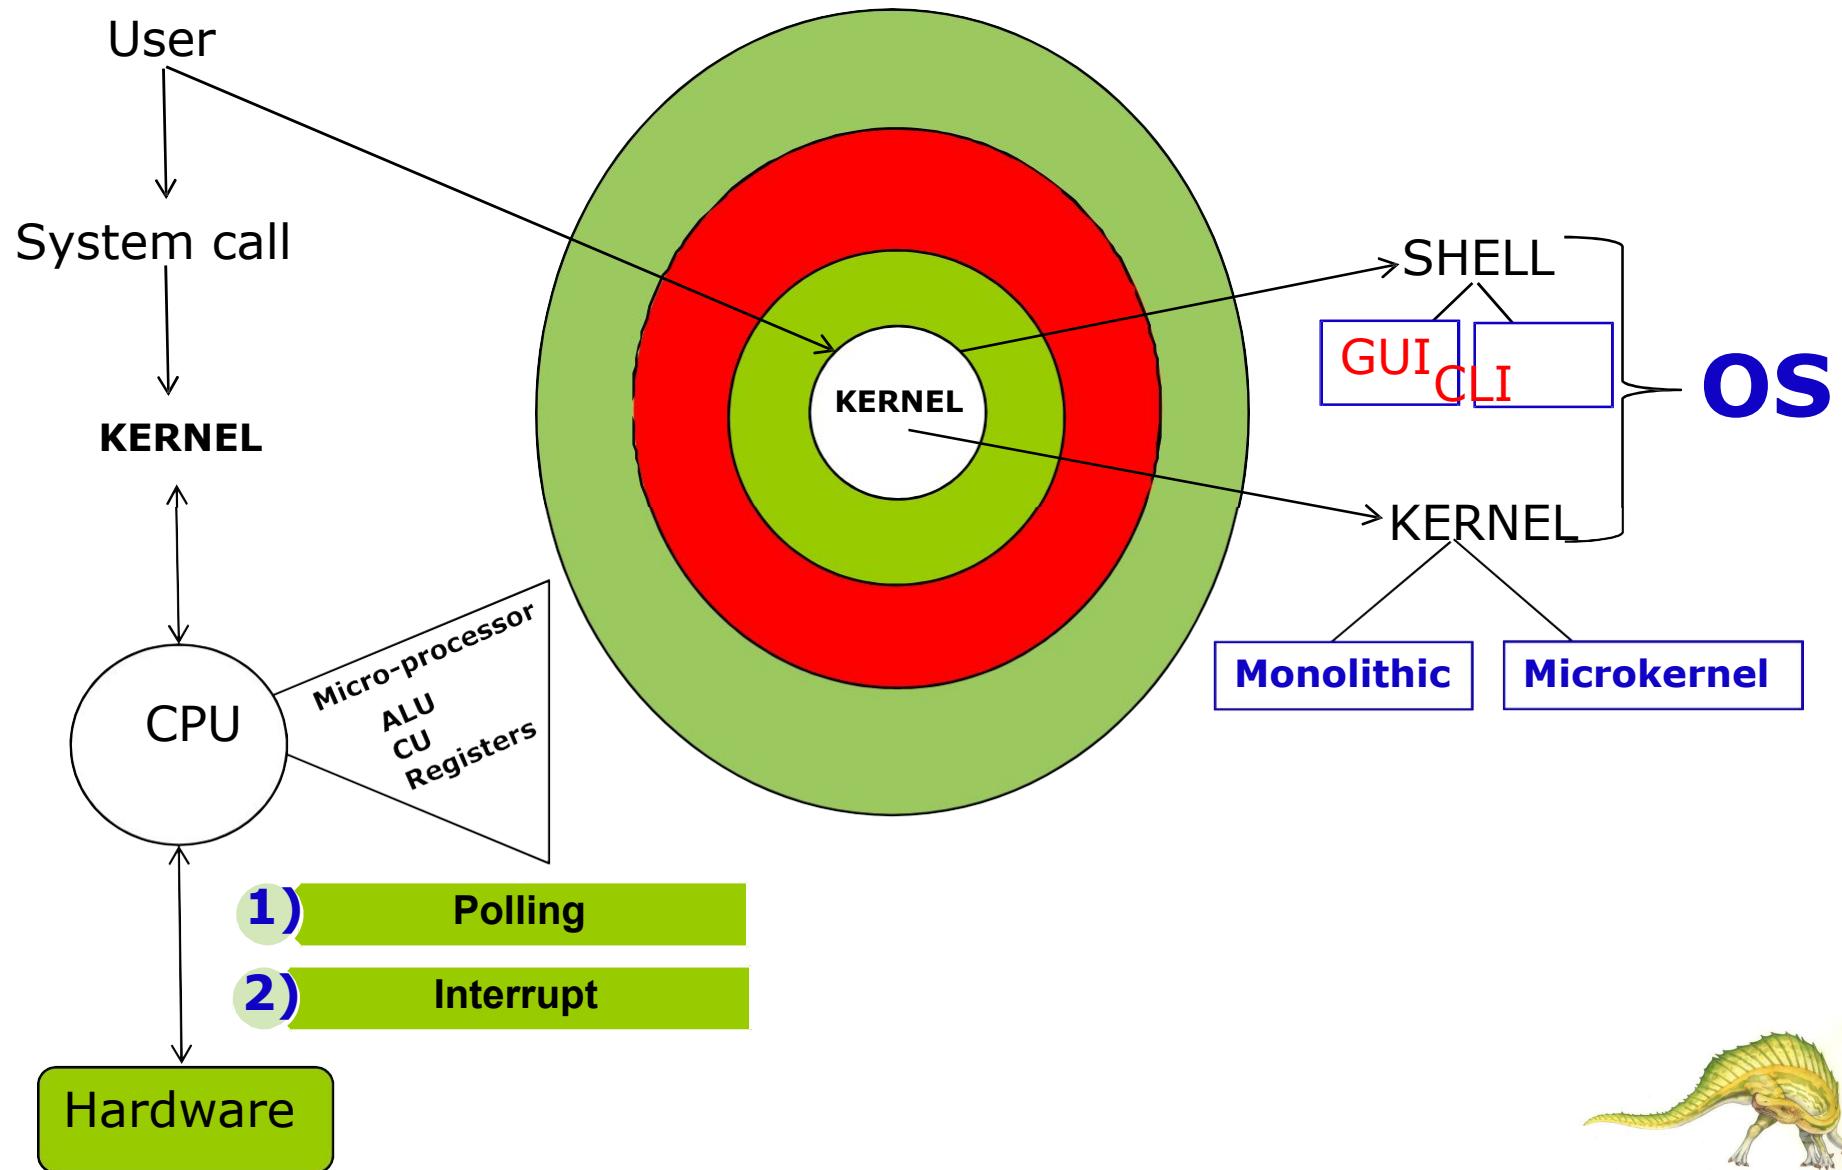
DMA – Direct Memory Access

- High Speed, I/O device





How does user communicate with OS





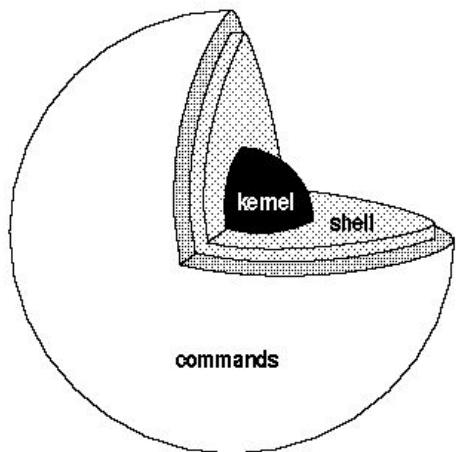
The Kernel and The Shell

An Operating system has two parts: **kernel** and **shell**

KERNEL (resides in main memory):

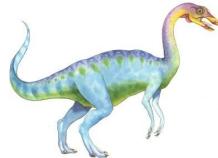
- is the **heart and soul** of an OS. Useful applications and utilities are added over the kernel, then the complete package becomes an OS.
- directly controls the computer hardware by performing **OS services using specific system calls or requests or hardware interrupts**.
- hides the hardware details from the user and application programs.

Example,



- **Linux is a kernel** as it does not include applications-file-system utilities, windowing systems and graphical desktops, system administrator commands, text editors, compilers etc.
- So, various companies add these kind of applications over Linux kernel and provide their operating system like **ubuntu, centOS, redHat** etc.





An Operating system has two parts: kernel and shell

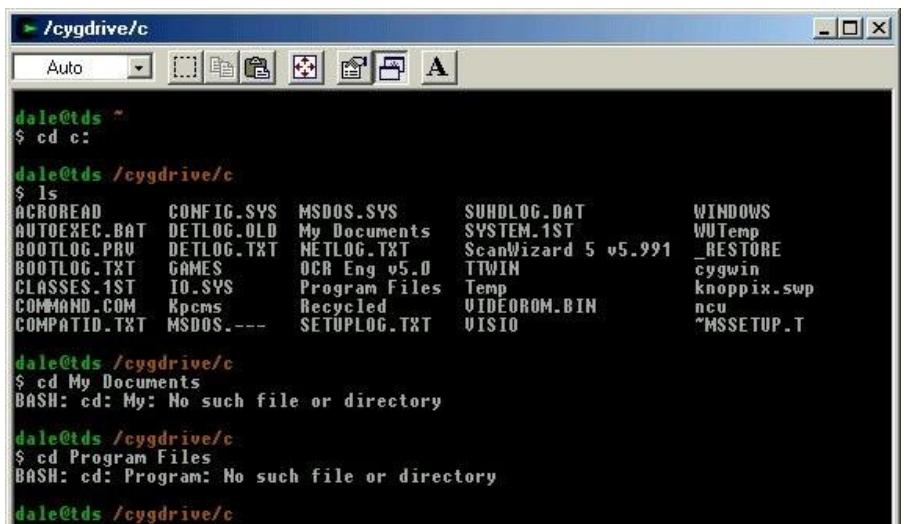
shell (command interpreter): Part of OS

- Serves as the interface between user and kernel.
- User enters commands via shell in order to use the kernel service.

Type of shells



Graphical User Interface (GUI)
Windows/Mac OS



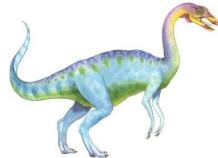
```
dale@tds ~
$ cd c:
dale@tds /cygdrive/c
$ ls
ACBREAD CONFIG.SYS MSDOS.SYS SUHDLOG.DAT WINDOWS
AUTODEC.BAT DETLOG.OLD My Documents SYSTEM.1ST WUTemp
BOOTLOG.PRV DETLOG.TXT NETLOG.TXT ScanWizard 5 v5.991 _RESTORE
BOOTLOG.TXT GAMES OCR Eng v5.0 TTWIN
CLASSES.1ST IO.SYS Program Files Temp
COMMAND.COM Kpcms Recycled VIDEOROM.BIN
COMPATID.TXT MSDOS.--- SETUPLOG.TXT VISTO
$MSSETUP.T

dale@tds /cygdrive/c
$ cd My Documents
BASH: cd: My: No such file or directory

dale@tds /cygdrive/c
$ cd Program Files
BASH: cd: Program: No such file or directory

dale@tds /cygdrive/c
```

Command line interface (CLI)
Dos/Linux



OS Protections

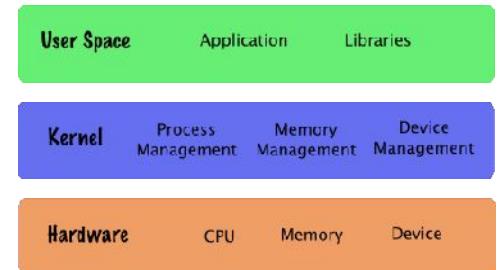




Transition from User to Kernel Mode

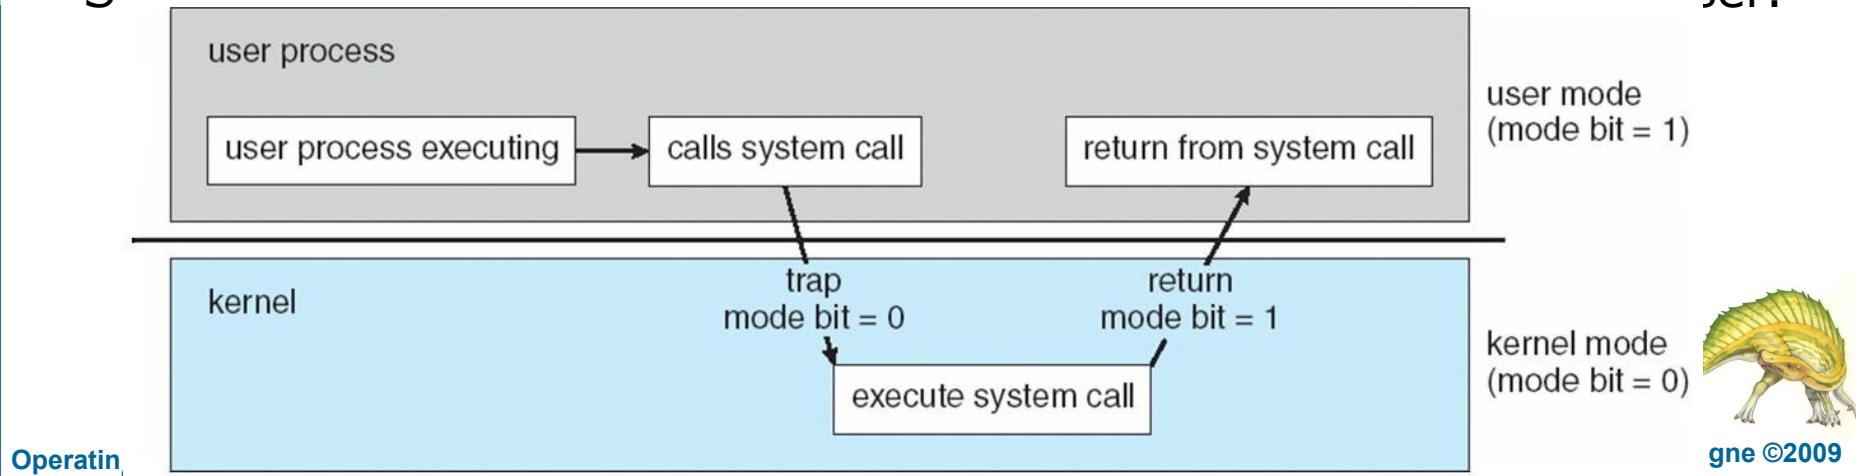
Dual-mode operation allows OS to protect itself and other system components.

- User mode
- Kernel/Supervisor/Monitor/System Mode.



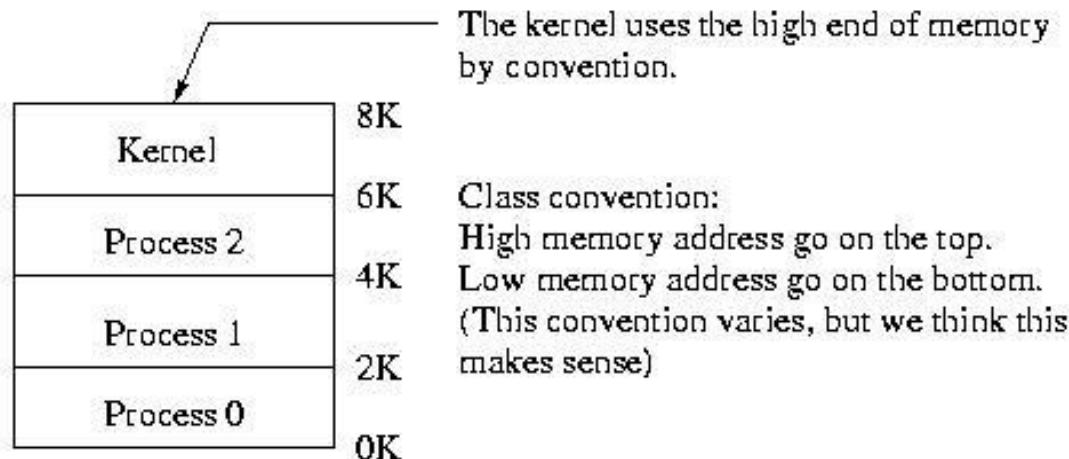
Mode bit @ Process Status Word (PSW)-Register

- Provides ability to distinguish when system is running user code or kernel code.
- Some instructions designated as **privileged**, only executable in kernel mode. Example: load/save from protected memory, initiate I/O etc.
- System call changes mode to kernel return from call resets it to user.





Simple Model for Memory Protection



To obtain true protection,
need more hardware support.

- The goal is to make sure that, for example, process 2 can't go outside of the 4K-6K region.
 - This requires hardware support -- we can't do it in software alone.
 - The simplest model is to add two registers: **BASE** & **LIMIT** (assigned by OS)
 - This is the simplest model, most computers are more complex.

In order to keep process 2 within its 2K box, we can do one of two things:

Option #1:

Set BASE = 4K

Set LIMIT=2K

A memory access is legal,

if $\text{BASE} \leq X < \text{BASE} + \text{LIMIT}$,
where X is the actual memory address

Option #2:

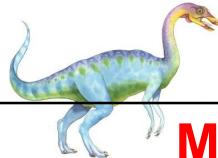
Set BASE =

LIMIT = 2K

more frequent scheme

A memory access is legal,

iff $0 \leq X < \text{LIMIT}$,
where X is the offset
address = $\text{BASE} + X$



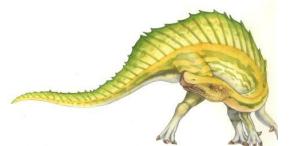
Kernel Types

Monolithic Kernel	Micro Kernel
All the parts of a kernel - Scheduler, file system, memory, drivers etc. are in one unit within the kernel.	Only the important parts are in kernel - IPC, basic scheduler, basic memory management, basic I/O primitives etc. - Others (including device drivers) are in user space.
Signals & sockets to ensure IPC	- Message Passing system to ensure IPC
Advantages - Faster processing	Advantages - Crash resistant - Smaller size - Portable - No need recompiling(for a new feature)
Disadvantages - Crash insecure - Porting Inflexibility - Add a new feature needs recompiling whole - Kernel size explosion	Disadvantages - slower processing due to additional message passing - 2%-4% slower
DOS, Unix, Linux, Symbian, Mac OS X, MINIX, Windows NT	





Lecture 2: OS Evaluation

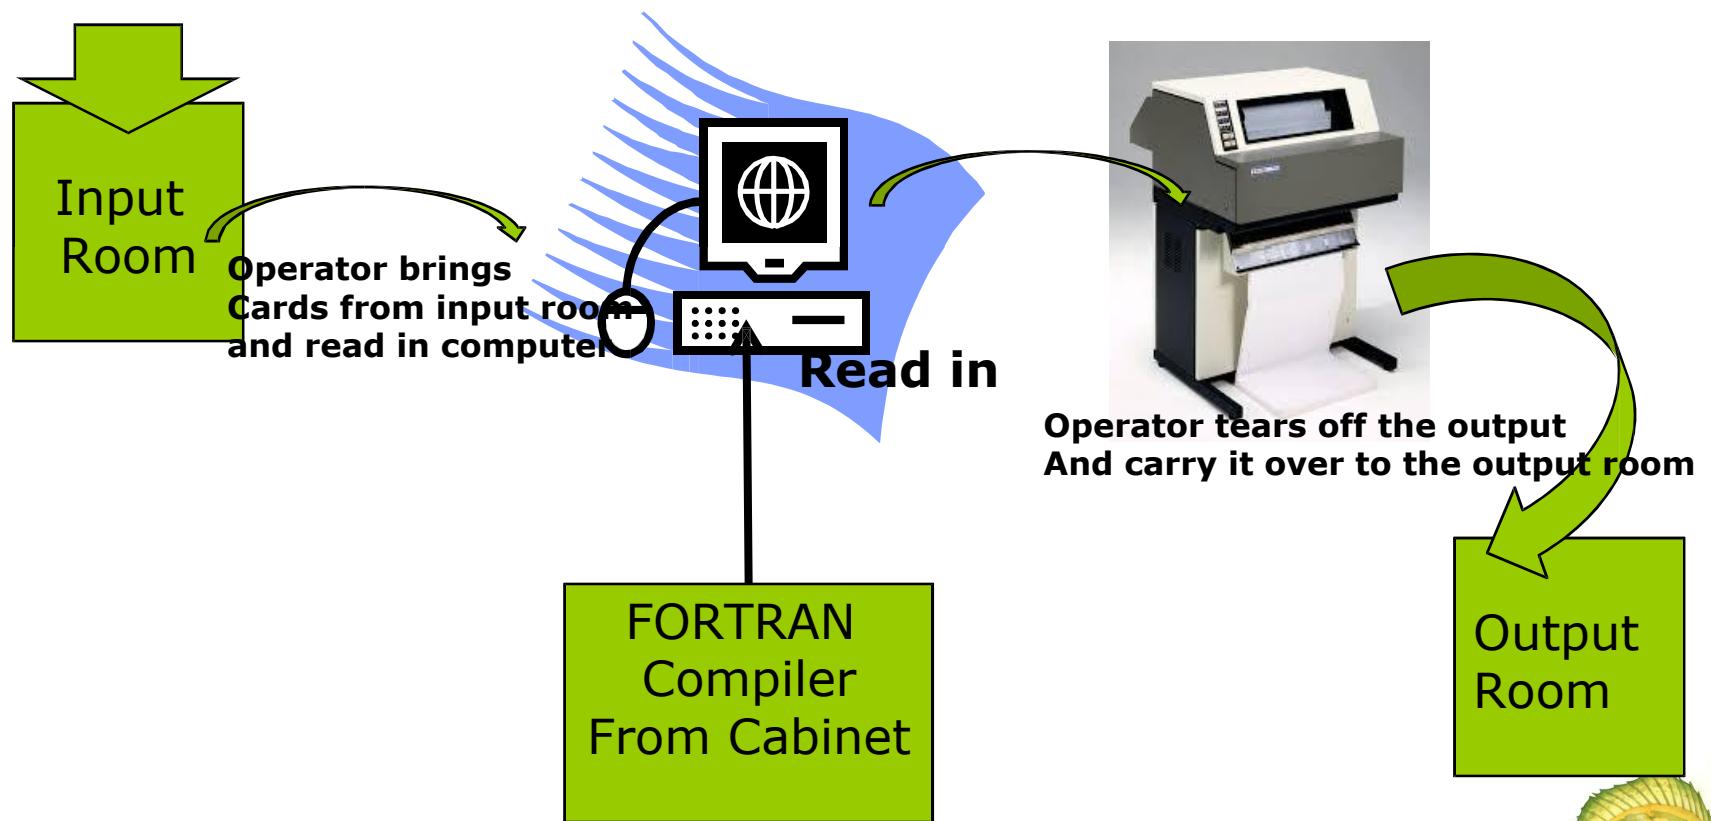




Computers were in Universities / Large companies

Programmers wrote programs (FORTRAN/ Assembly) on paper.

Punched the program into CARDS



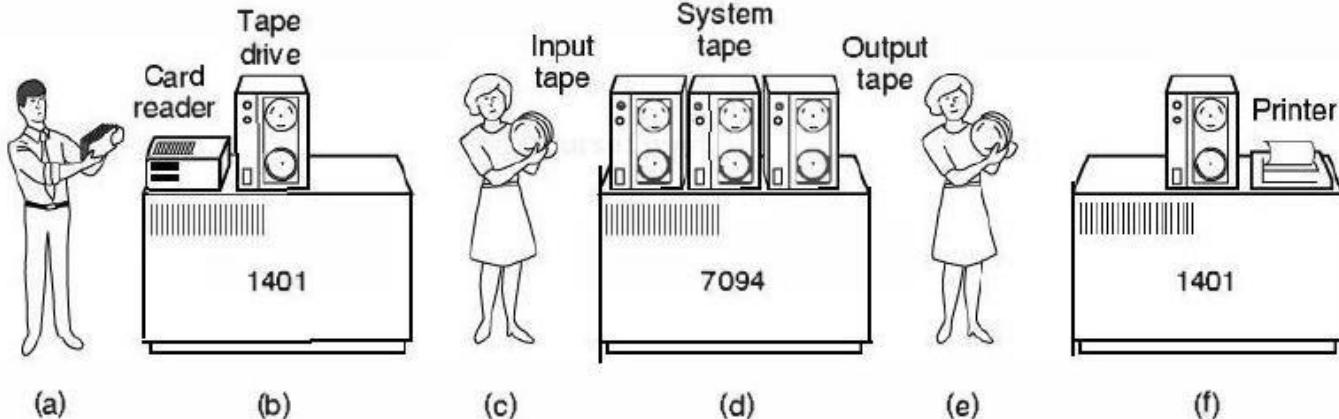
Wasting time while operators were walking around the machine room



Supervisor/Operator Control

CR ----> MT MT ----> CPU CPU ----> LP

IBM 7094



An early batch system. (a) Programmers bring cards to 1401. (b) 1401 reads batch of jobs onto tape. (c) Operator carries input tape to 7094. (d) 7094 does computing. (e) Operator carries output tape to 1401. (f) 1401 prints output.

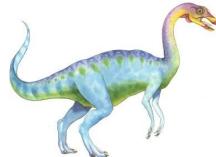
- Problems
 - Long turnaround time - up to 2 DAYS!!!
 - Debugging...
 - Low CPU utilization
 - I/O and CPU could not overlap;
 - slow mechanical devices.



From John Ousterhout slides

30





BATCH Processing (One by one) Similar jobs

Process 1		Process 2		Process 3		
I1		I2		I3		
	C1		C2		C3	
	O1		O2		O3	
2	4	2	2	3	2	7
						26



I= Input
C=Computation
O=Output

OS resides in memory and the task was simple
- transfer control from one job to another

Problem: Low CPU utilization;
- During I/O CPU is idle. WHY?



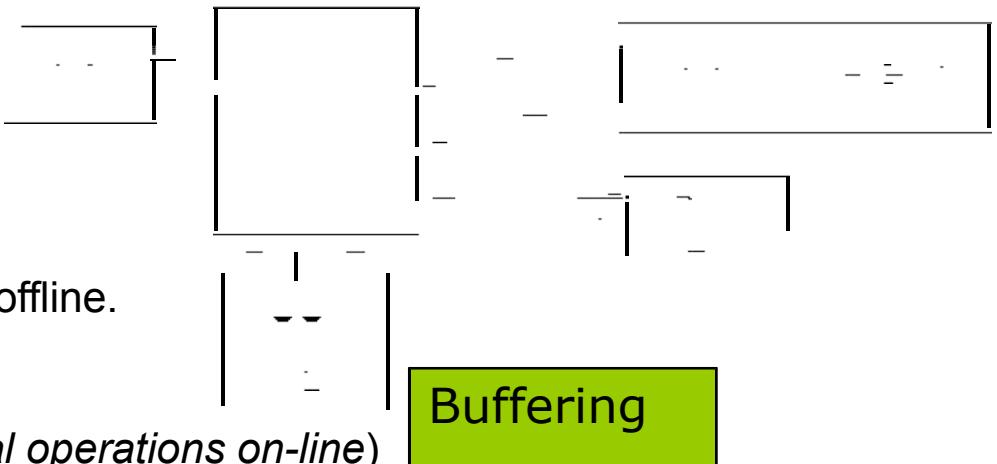


Batch Systems - Issues

- **Solutions to speed up I/O:**

- Offline Processing (Previous)

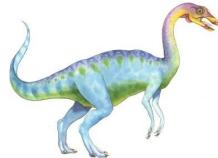
- load jobs into memory from tapes,
 - card reading and line printing are done offline.



- Online Spooling (*simultaneous peripheral operations on-line*)

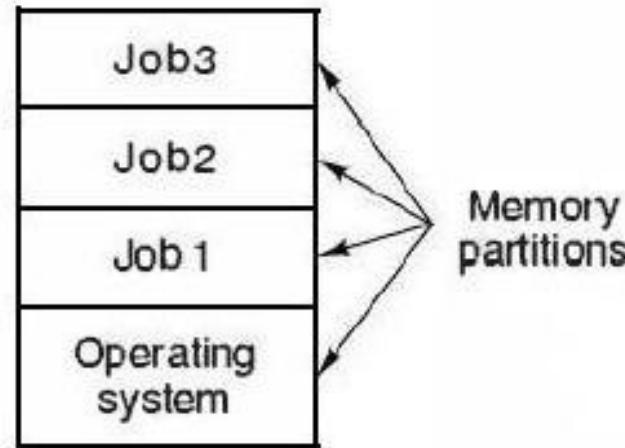
- A type of buffering, to disk (present)
 - Random access device, unlike tape
 - Use disk as large storage for reading as many input files as possible and storing output files until output devices are ready to accept them.
 - Allows overlap - I/O of one job with computation of another.
 - Introduces : job pool (Queues-job queue, ready queue)
 - allows OS to choose next job to run so as to increase CPU utilization





Multiprogramming

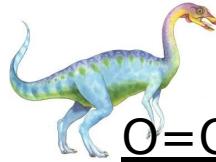
- Use interrupts to run multiple programs simultaneously
 - When a program performs I/O,
 - instead of polling, execute another program till interrupt is received.
 - Efficiency



A multiprogramming system with three jobs in memory.

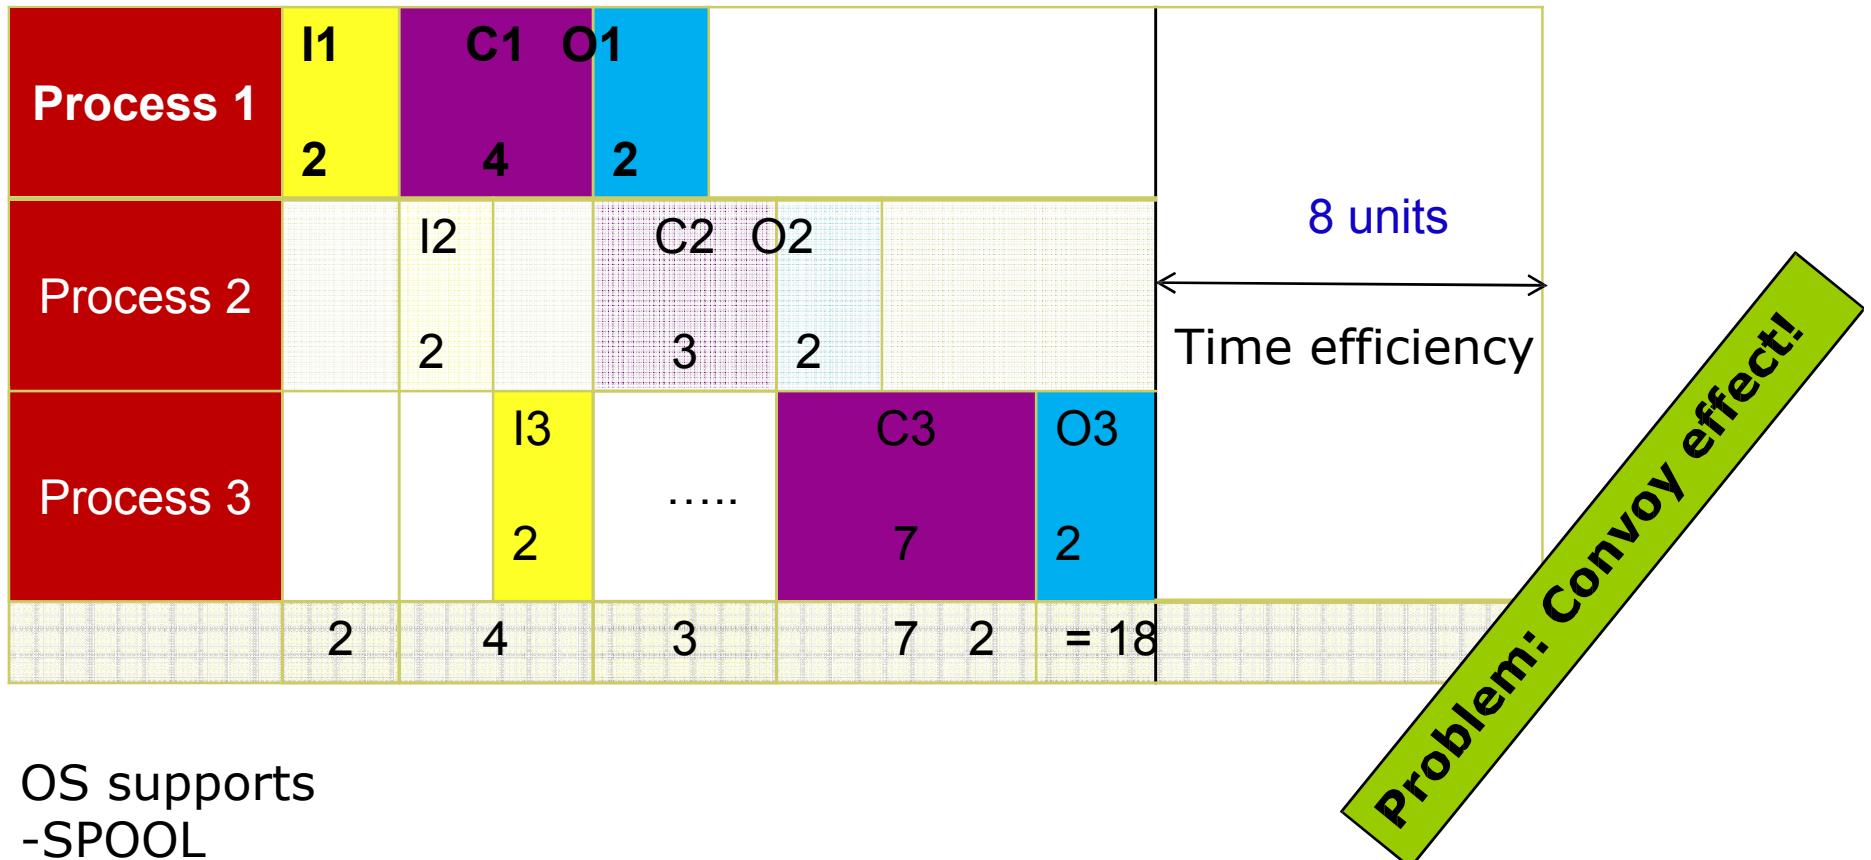
- Requires:
 - **secure memory [will learn @ final] and I/O for each program.**
 - **intervention, if program loops indefinitely.**
 - CPU scheduling to choose the next job to run.[will learn later]



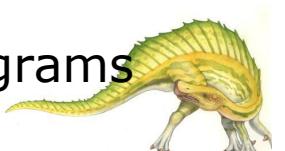


Multiprogrammed Processing

I= Input
C=Computation
O=Output



- OS supports
- SPOOL
- Overlapping I/O and CPU [I/O of one job CPU of another job]
- Interleaving – mixed of CPU and I/O operations of several programs and alternating between them





Timesharing

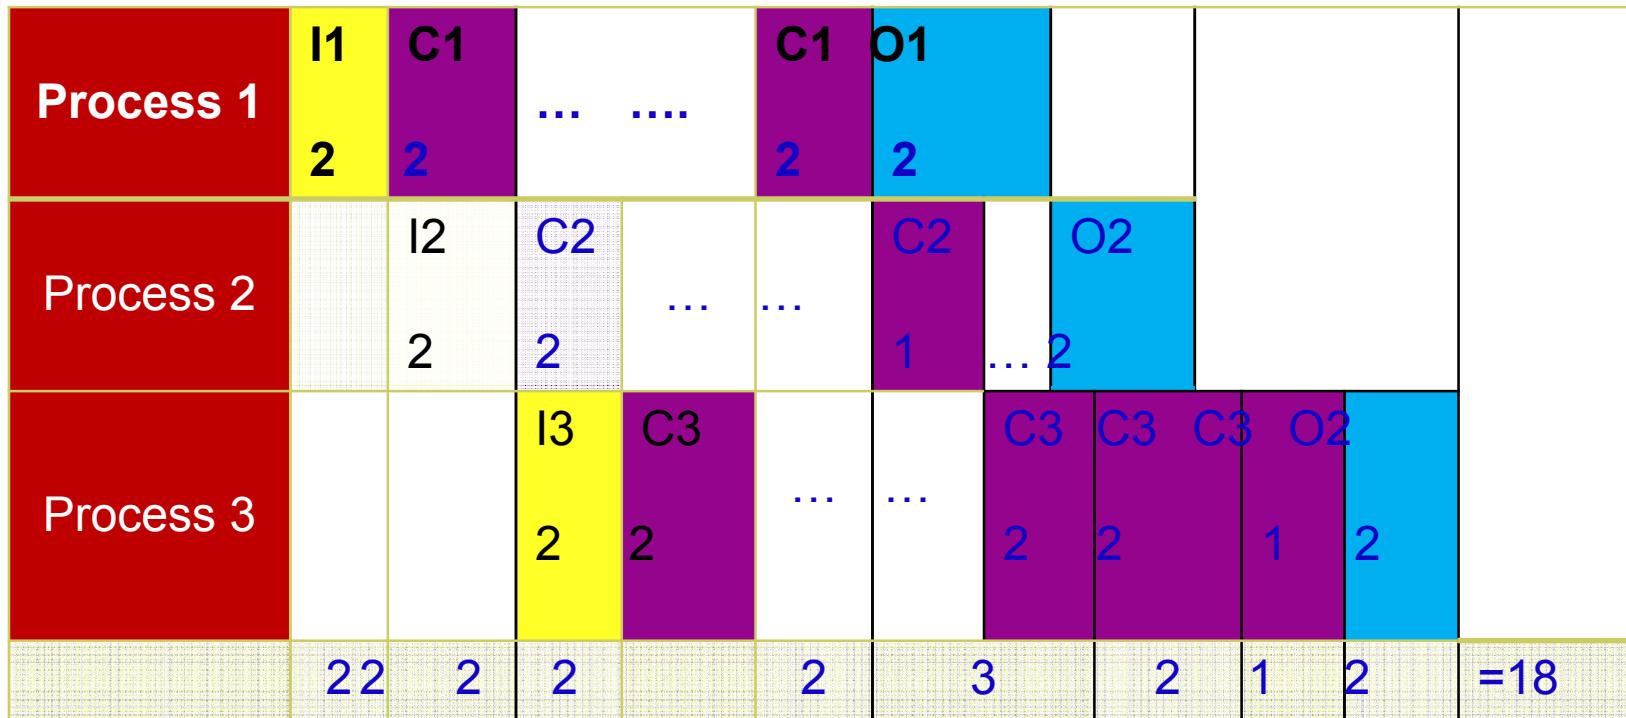
- A variant of multiprogramming.
- Supports multiple users at the same time
- Programs queued for execution in FIFO order.
- A technique to protect CPU.

- **Like multiprogramming, but timer device interrupts** after a quantum (timeslice).
 - Interrupted program is returned to end of FIFO
 - Next program is taken from head of FIFO





Time Shared Processing



Concept of Virtual Processors
Time Slice = 2



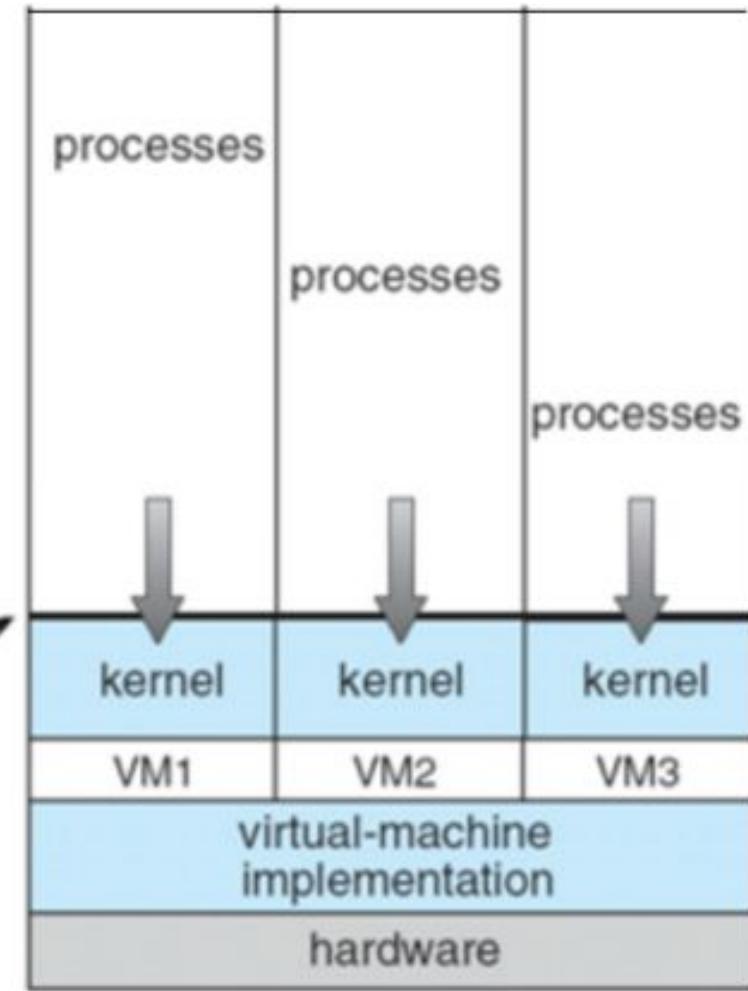


Virtual Machines



(a)

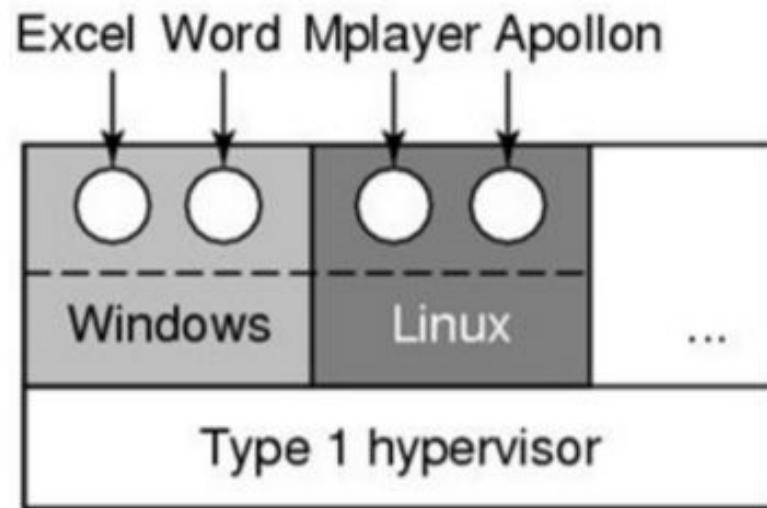
(a) Nonvirtual machine



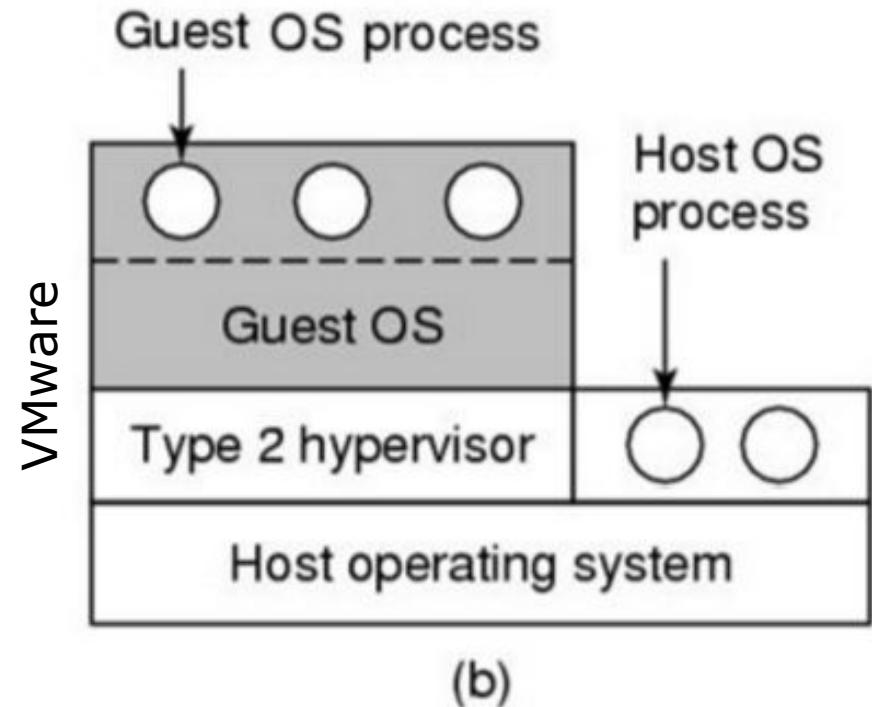
(b)

(b) virtual machine





(a)



(b)

(a) A type 1 hypervisor. (b) A type 2 hypervisor.





Reading Materials

- Galvin : 1.1, 1.2, 2.1, 2.2, 2.5, 3.3, 3.5, 3.6
- Galvin: 13, 13.2.1, 13.2.2



End of Lecture 1-3

