




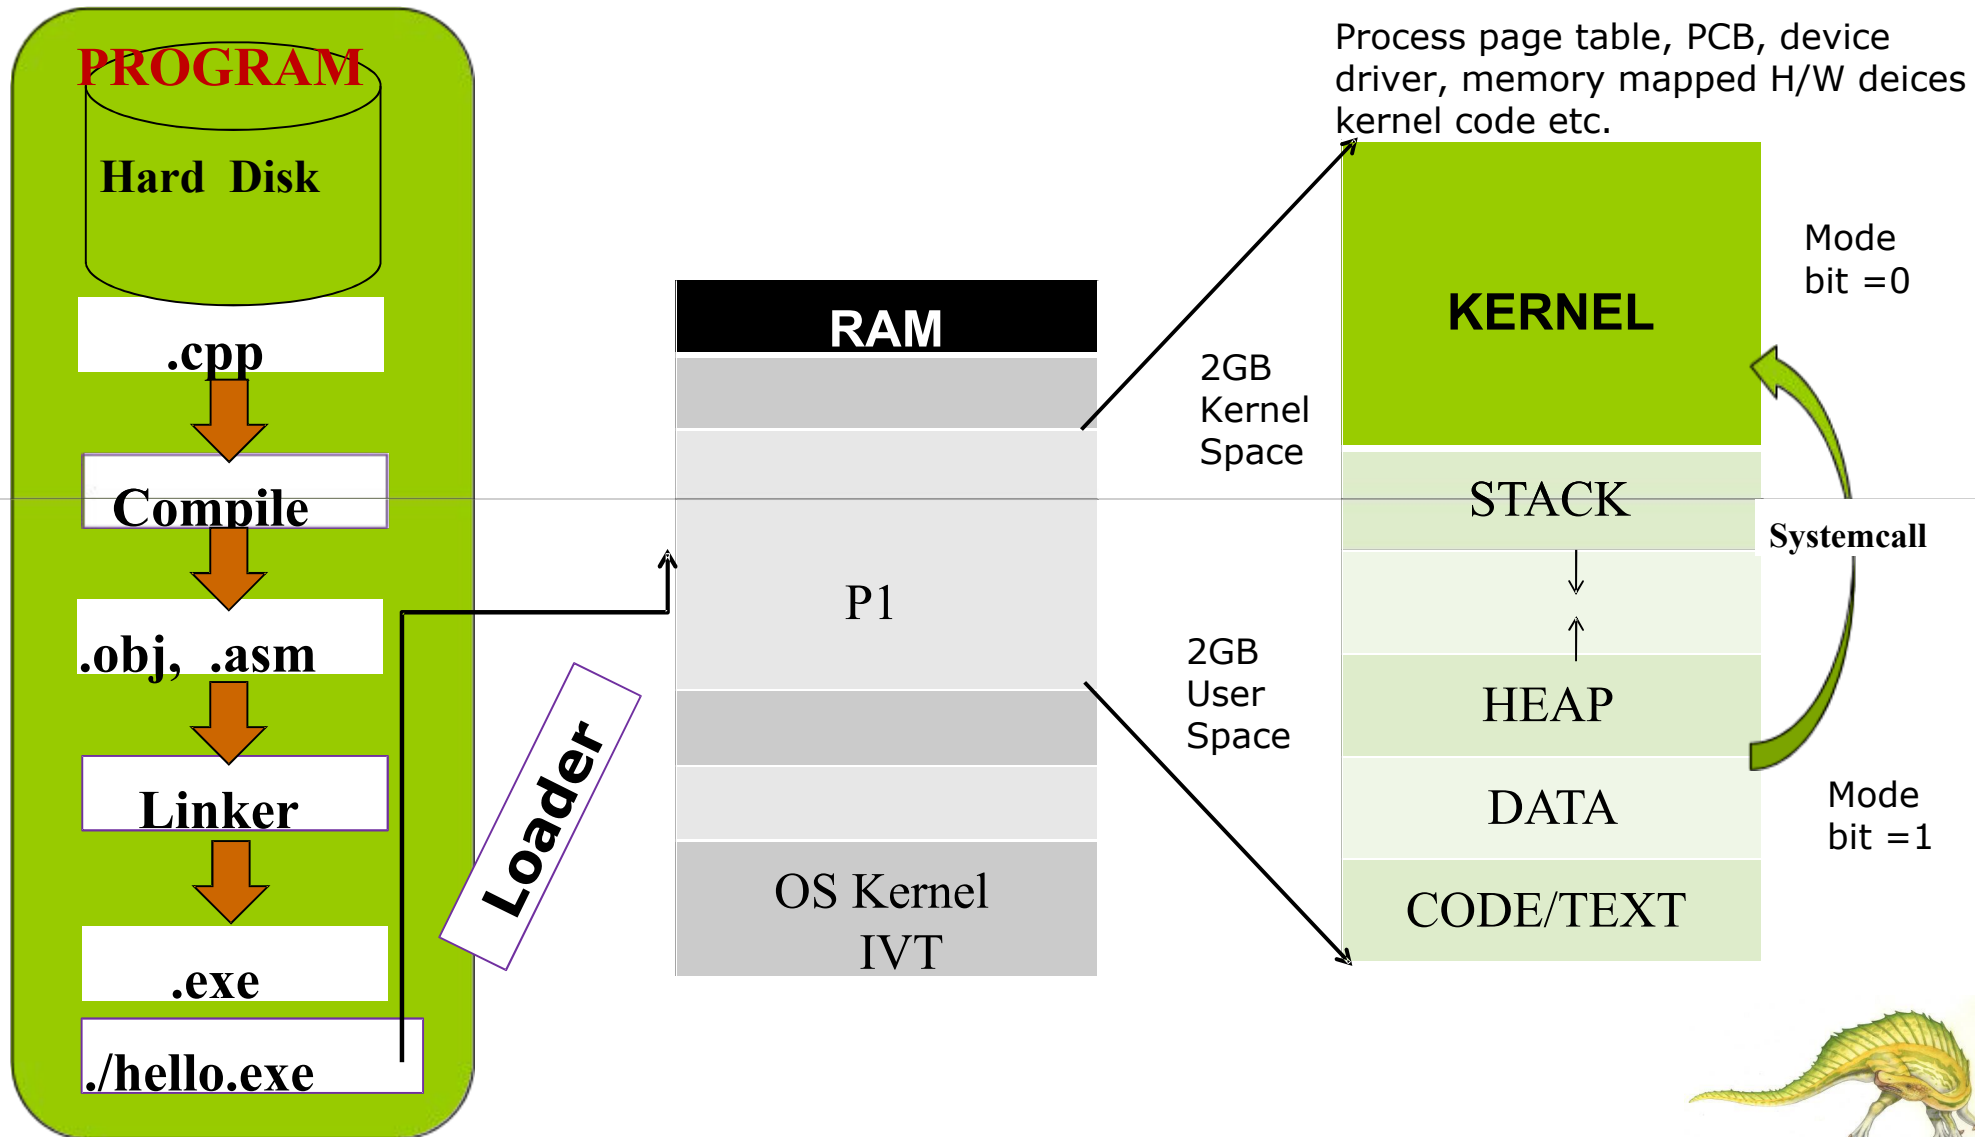
Lecture 4, 5


Process Management





Program to Process





Memory Organization for an Executed Program

When a program is loaded into memory, user space is organized into **four regions of memory**, called segments:

text segment, data segment, stack segment & heap segment

Text segment (or code segment)

where the compiled **code of the program** itself resides.

Data segment (Data & BSS)

data area contains

global or static or external variables that are **initialized**.

BSS contains

global or static or external variables that are **uninitialized**.

Pointer variable `int *arr` ; declared in global then in data else in stack


Heap segment

dynamically allocated variables are allocated in here.

it is managed by `malloc` and `free`.

Stack segment

- contains the program **stack**, a **LIFO** structure.
- `$sp` register point to the top of the stack.
- memory is allocated for **automatic variables (local variables)** within functions scope.





The Process-Executable Program

- We write a program in e.g., Java.
- A compiler turns that program into an instruction list.
- The CPU interprets the instruction list (which is more a graph of basic blocks).


```
void X (int b) {  
    if (b == 1) {  
        ...  
    }  
}  
  
int main() {  
    int a = 2;  
    X(a);  
}
```

Program counter
next instruction address



Process in Memory

max



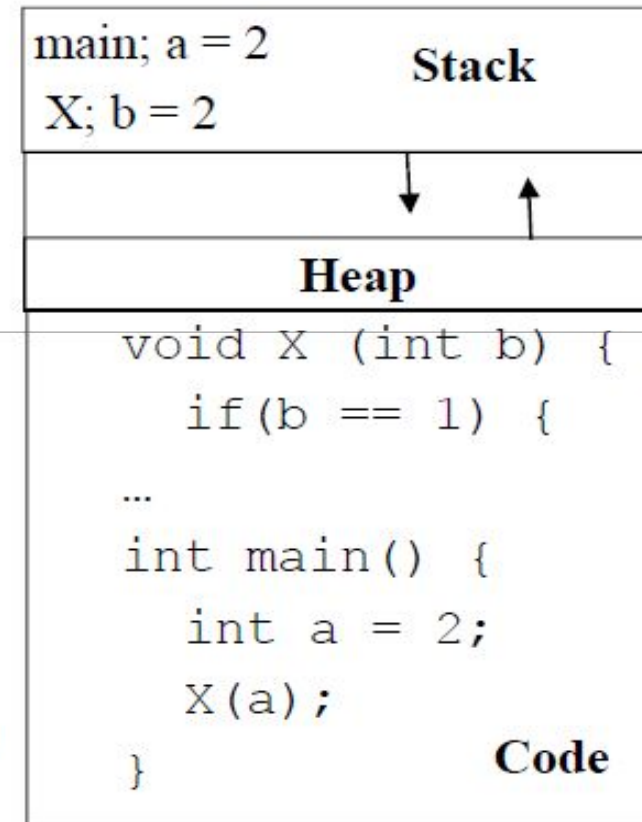
- Program to process.

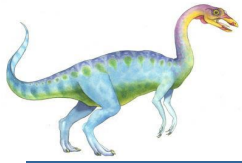
- What you wrote

```
void X (int b) {
    if (b == 1) {
        ...
    }
int main() {
    int a = 2;
    X(a);
}
```

- What must the OS track for a process?

- What is in memory.





Process Concept @ OS

Textbook uses the terms job and process almost interchangeably.

A process is - **Execution of an individual program.**

Each time a process is created,

OS must create a complete independent **address space (base, limit)**
(i.e., processes do not share their heap or stack data)

RAM
P1
P2
P3

Represents by a Data Structure to OS
Called **Process Control Block- PCB**





Process Control Block (PCB)

OS maintains a **process table** to keep track of the active processes

Information for each process:

Program counter

Program id, user id, group id

Program status word

CPU register values

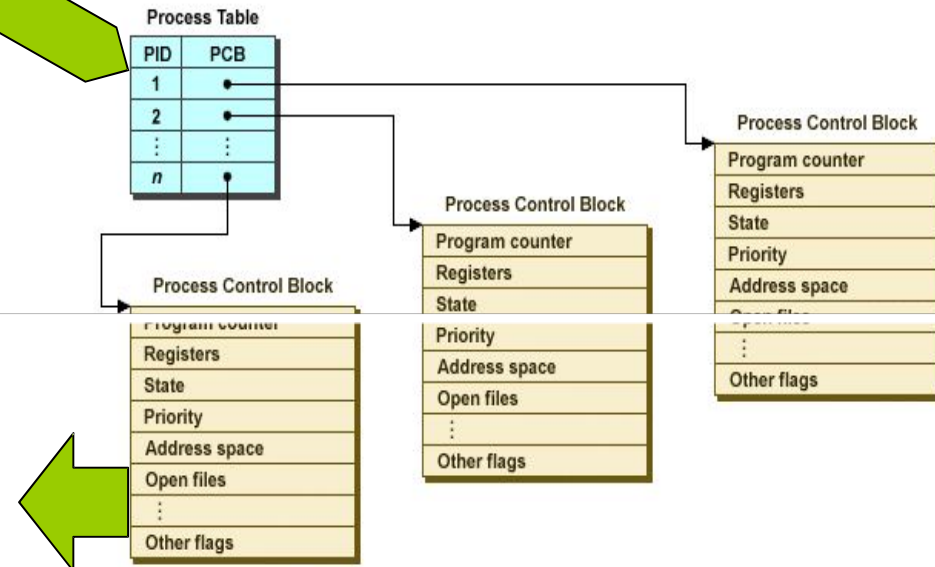
CPU Scheduling-process priority, pointer to scheduling queue

Memory maps-base/limit register, page table, segment table

Stack pointer

I/O status Information-allocated I/O devices, list of Open files

Accounting information, etc.-amount of CPU & real time used,
time limits, account numbers, job/process number



Stay in kernel
(Main Memory)




Two-State Process Model

Process may be in one of two states

Running

Not-running

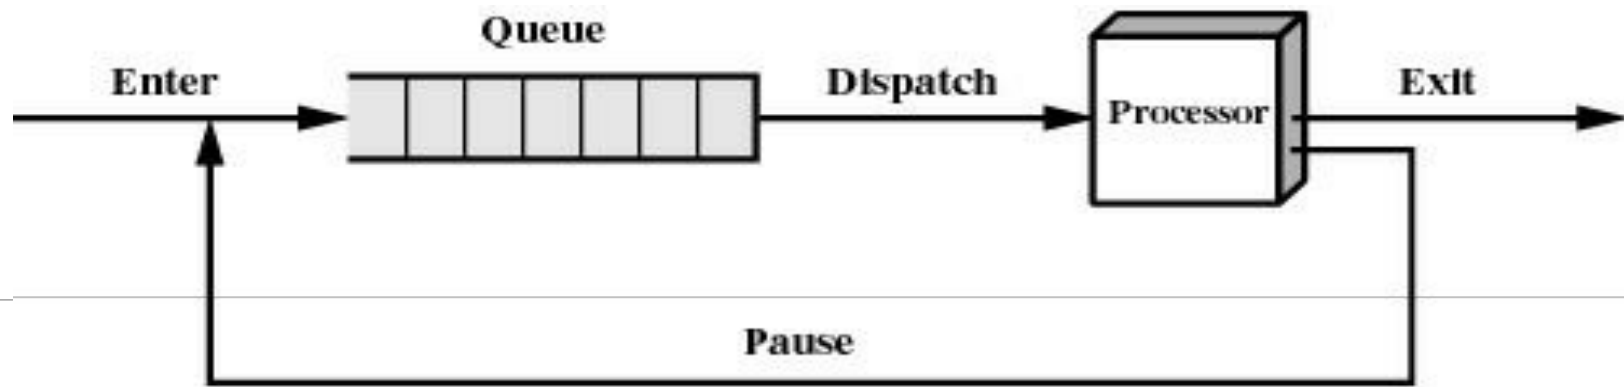


(a) State transition diagram



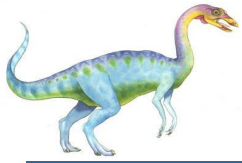


Not-Running Process in a Queue



(b) Queuing diagram





Five States Process Model

As a process executes, it changes *state*

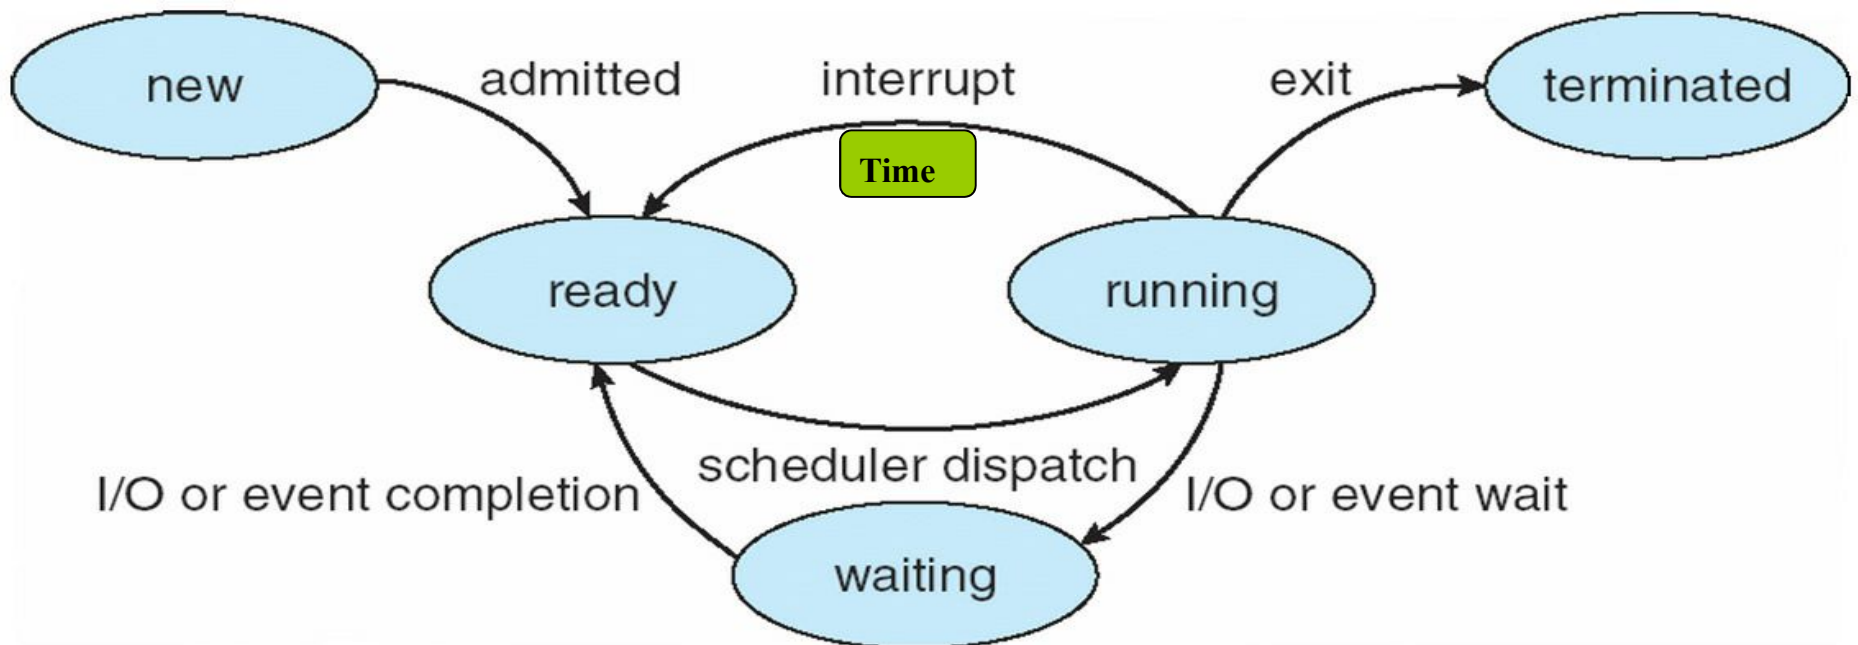
new: The process is being created

running: Instructions are being executed

waiting: The process is waiting for some event to occur

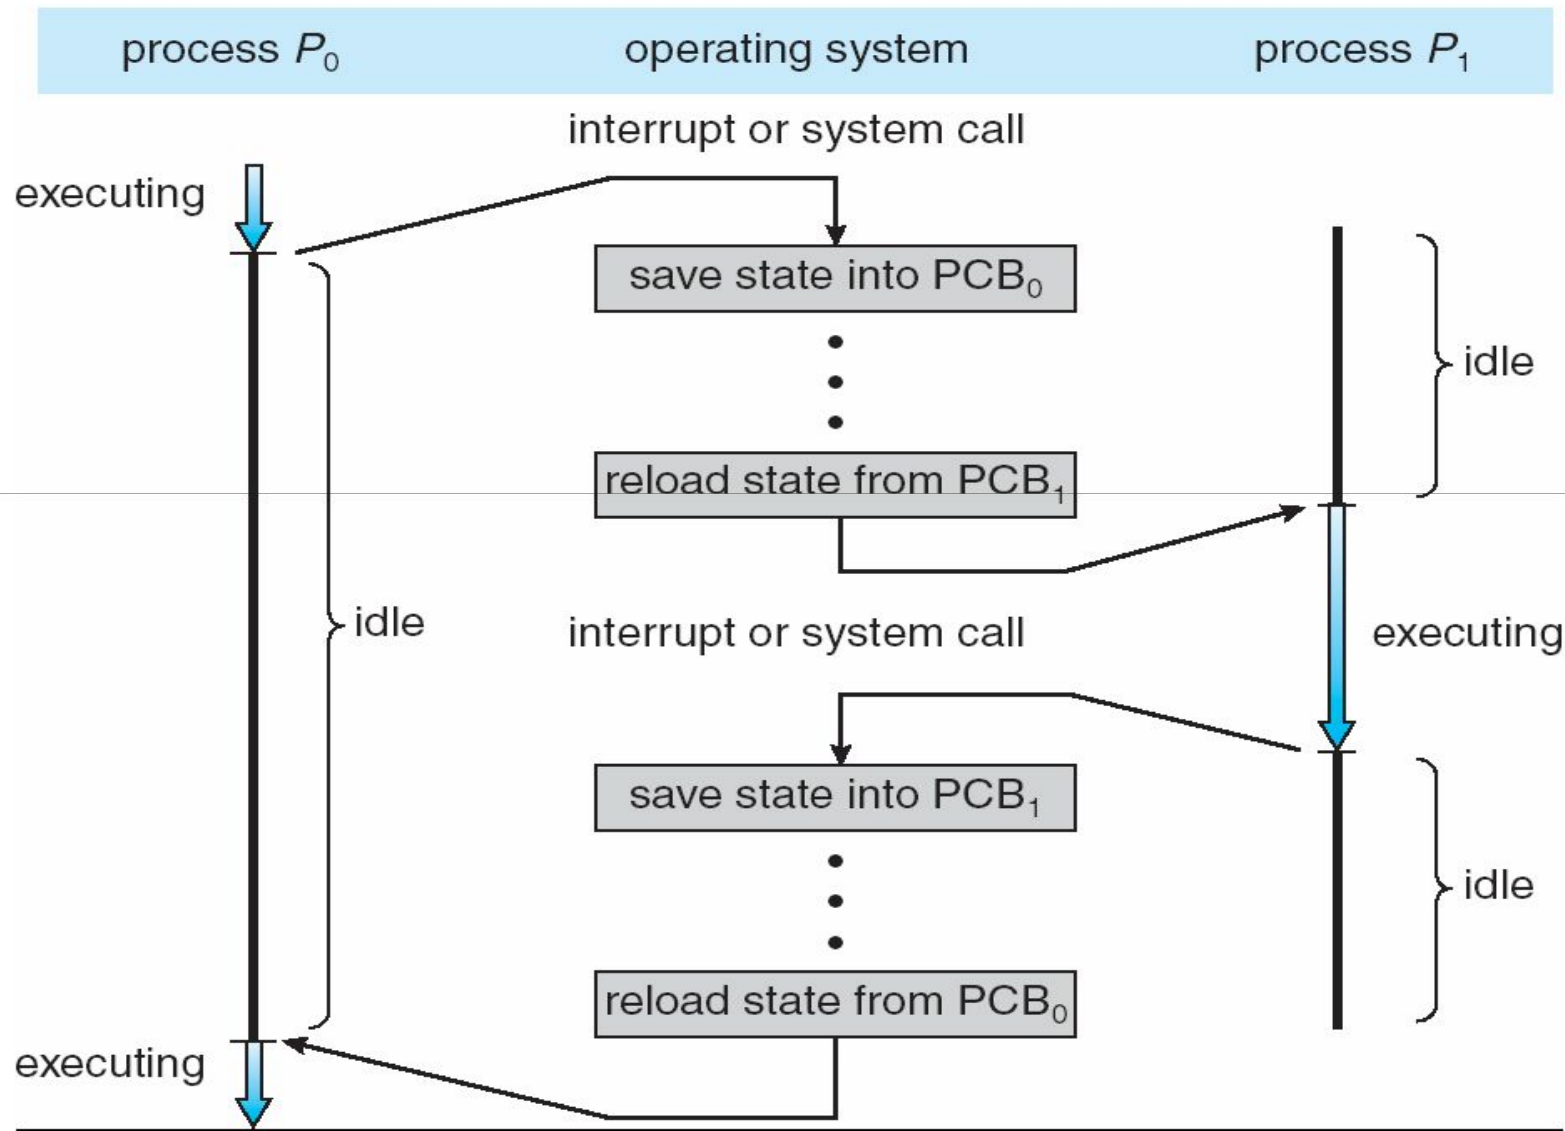
ready: The process is waiting to be assigned to a processor

terminated: The process has finished execution






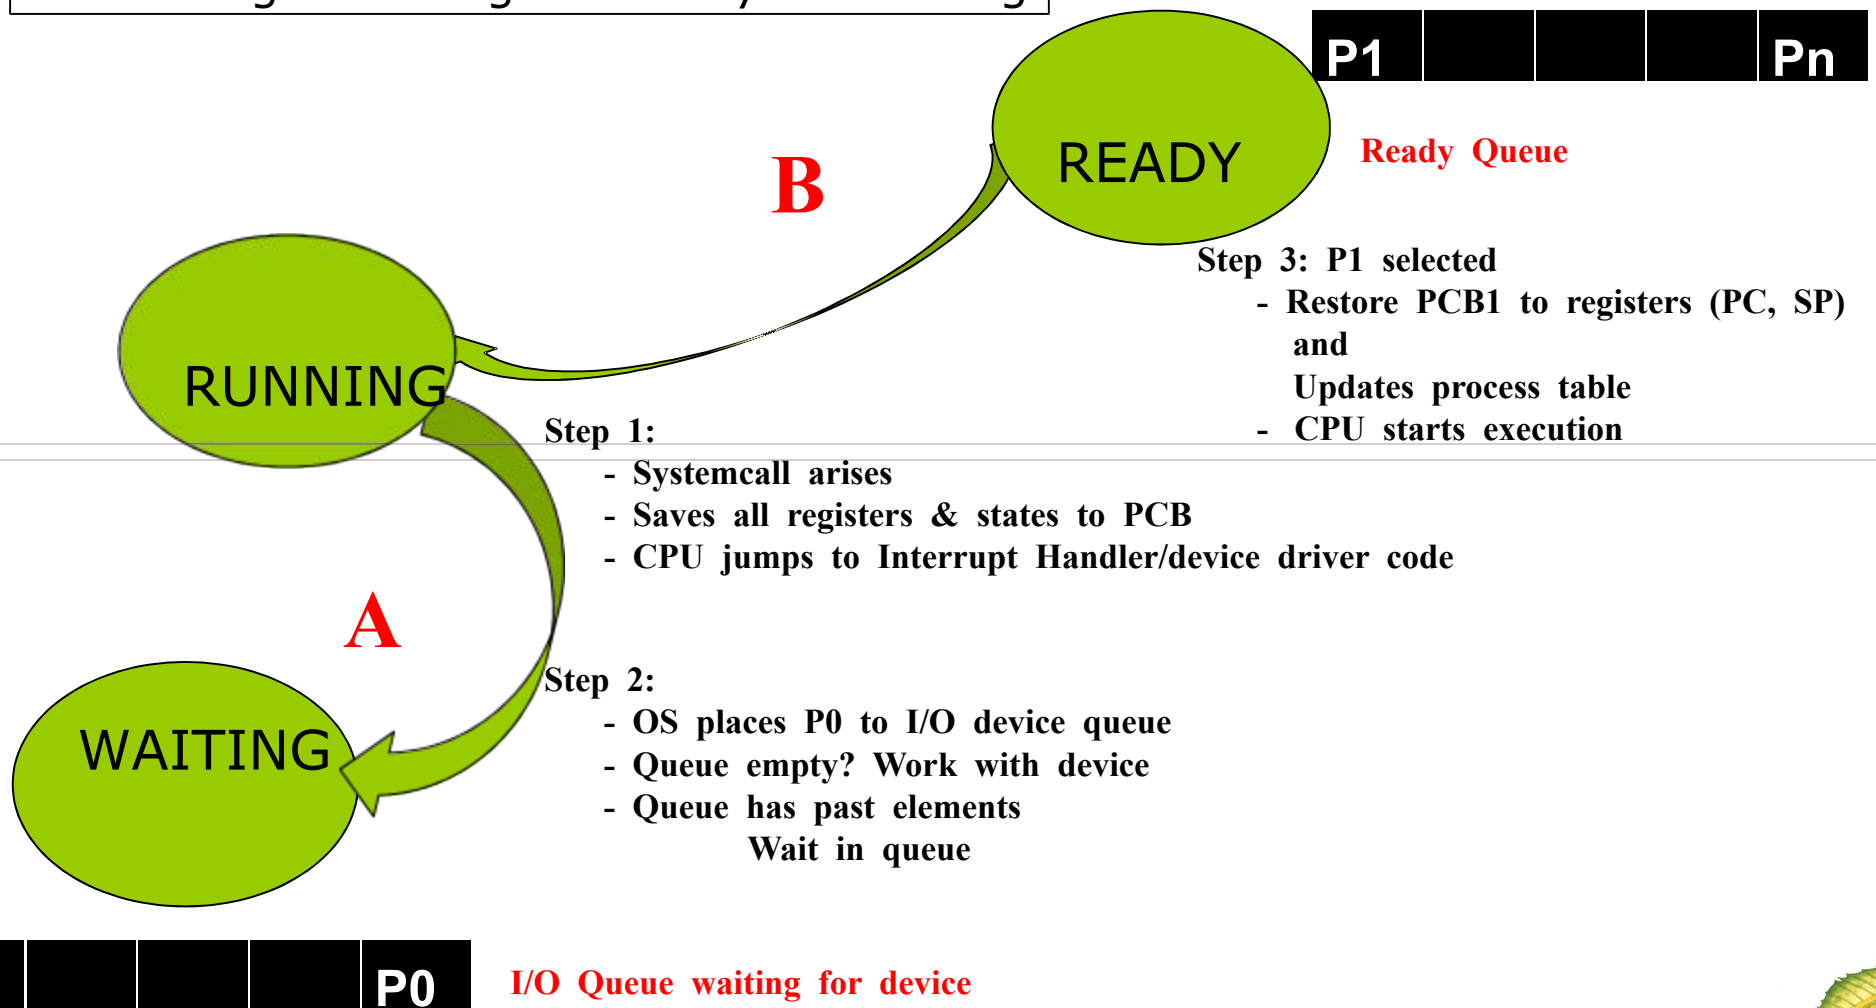
CPU Switch From Process to Process





Context Switch

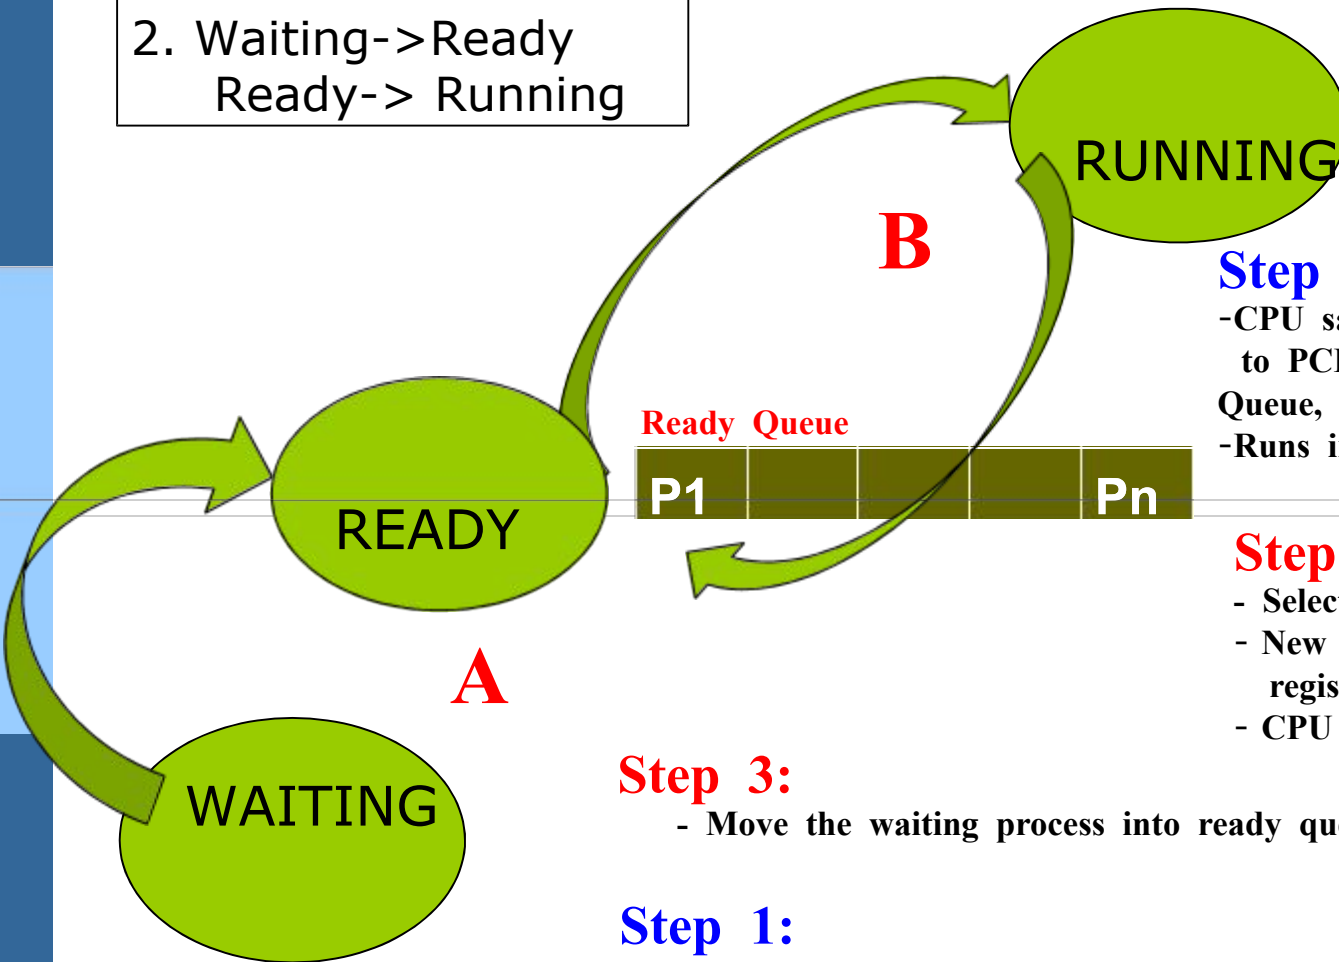
1. Running->Waiting && Ready-> Running





Context Switch

2. Waiting->Ready
Ready-> Running



Step 2:

- CPU saves running process Registers to PCB (and put the process back to Ready Queue, if required)
- Runs interrupt handler routine

Step 4

- Select the new process from Ready Queue
- New process PCB is transformed to register
- CPU starts execution


Step 3:

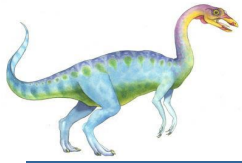
- Move the waiting process into ready queue (P0)

Step 1:

- Device finishes the job (P0)
- Create interrupt to CPU, OS executes interrupt handler
- Device is assigned next job

I/O Queue @ Device status table





Process Scheduling

AIM: Maximize CPU use, quickly switch processes onto CPU for time sharing

Process scheduler (is a process) selects process among available for next execution on CPU

Maintains **scheduling queues** of processes


Job queue – set of all processes in the system (HD-Pool)

Ready queue – set of all processes residing in main memory, ready and waiting to execute


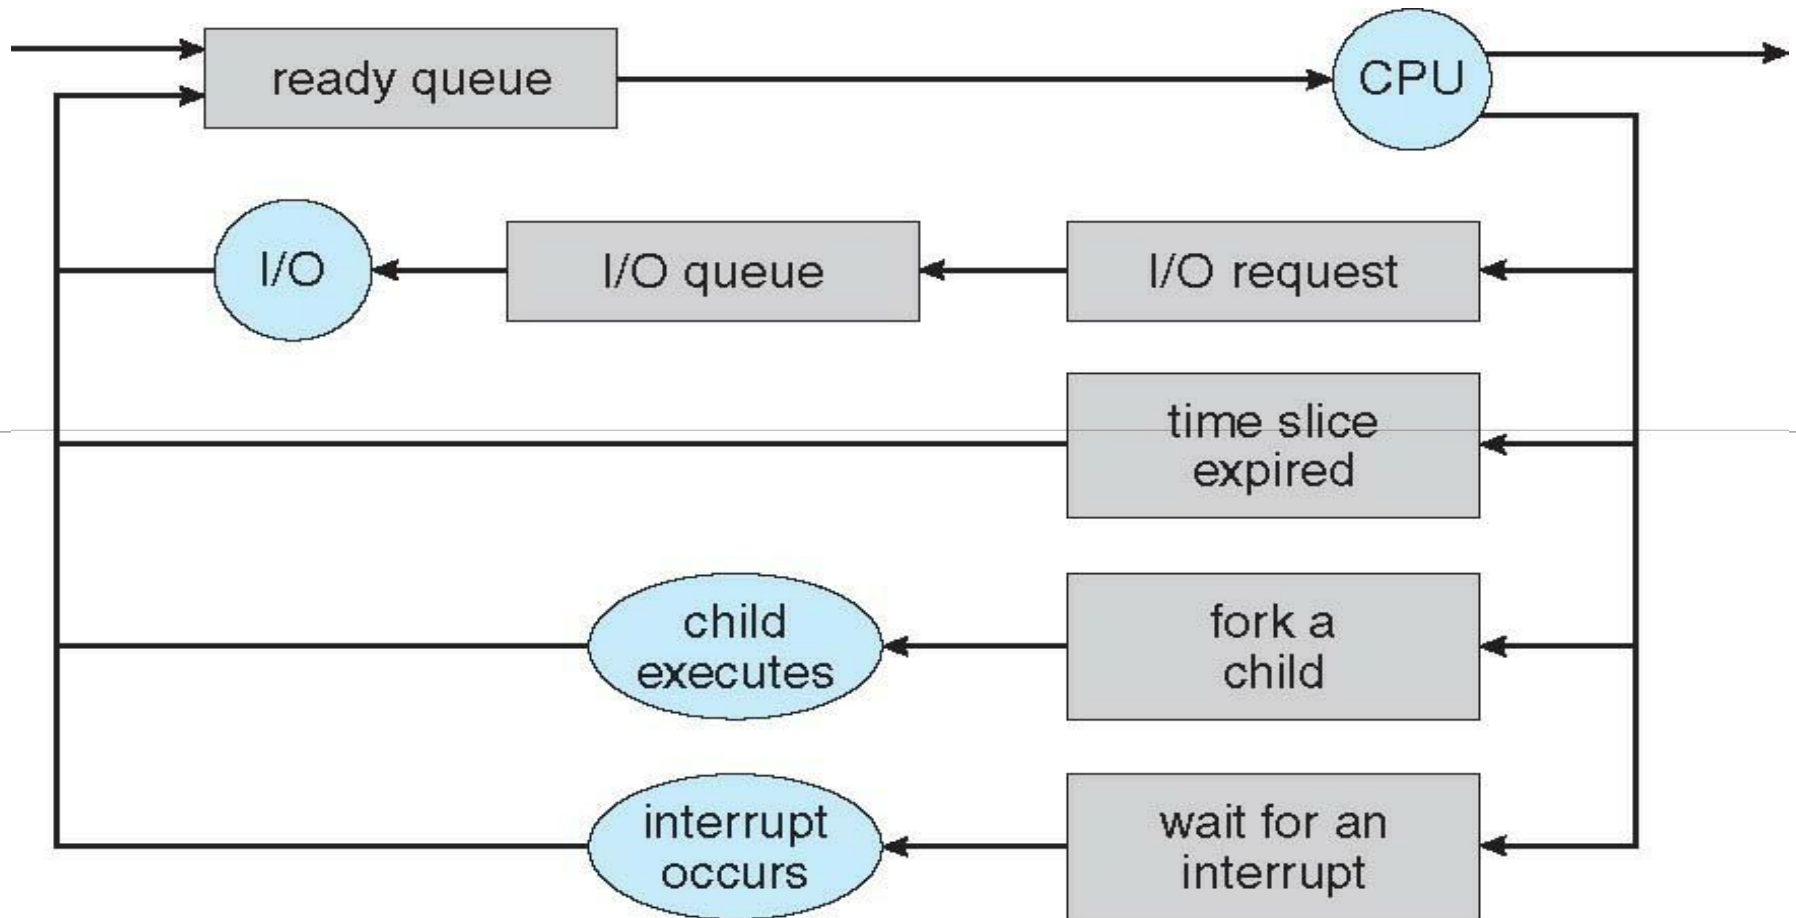
Device queues – set of processes waiting for an I/O device

Processes migrate among the various queues





Representation of Process Scheduling





Schedulers

Long-term scheduler/Job scheduler

which processes should be brought into the ready queue

invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)

controls the *degree of multiprogramming*

Short-term scheduler/CPU scheduler


which process should be executed next and allocates CPU

invoked very frequently (milliseconds) \Rightarrow (must be fast)

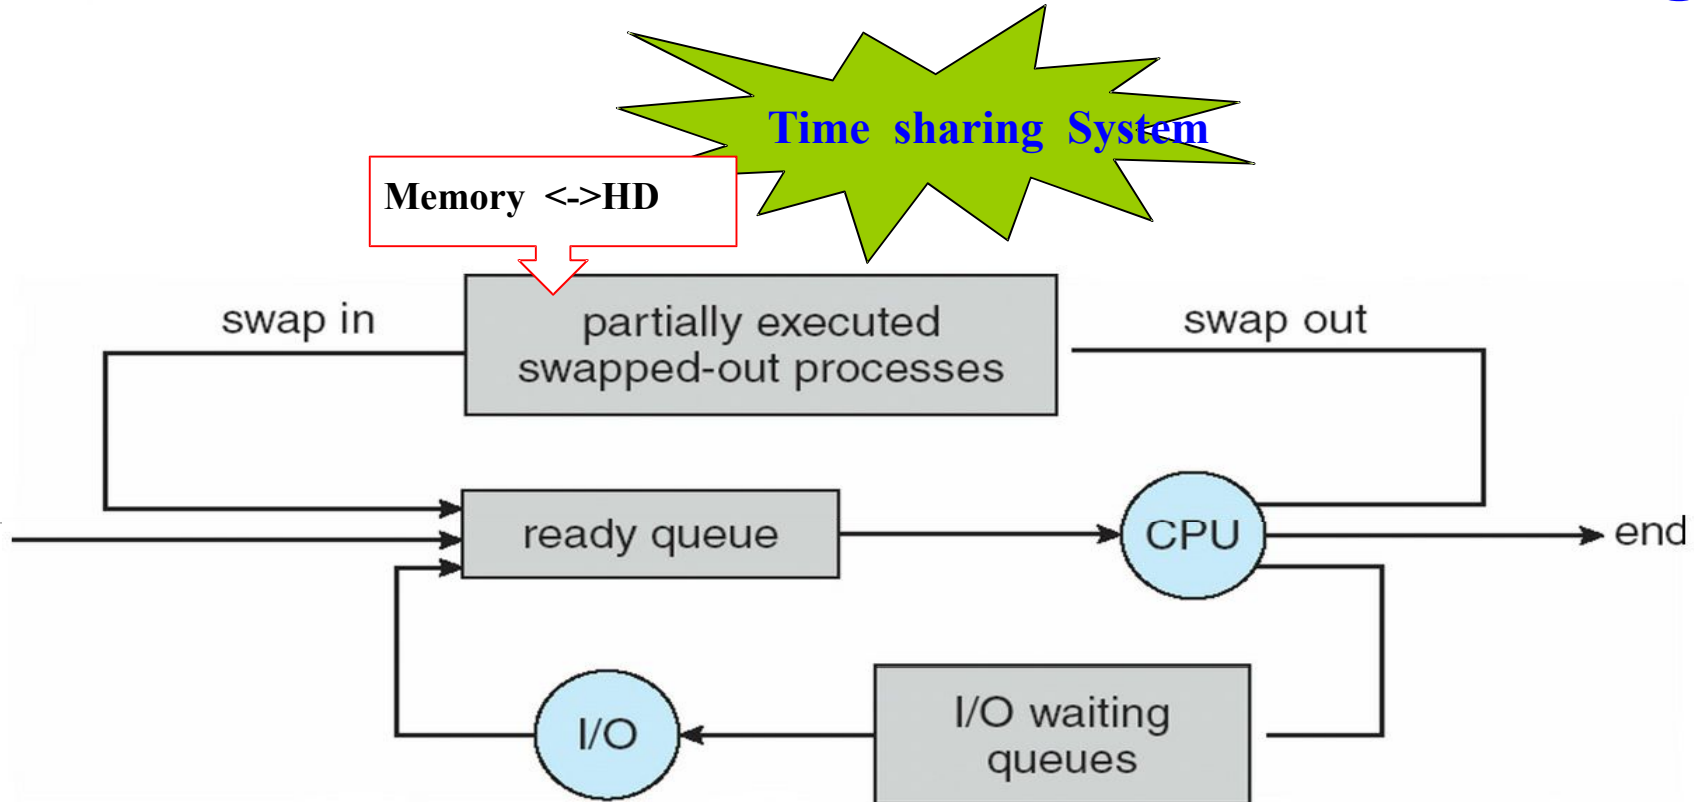
Balanced

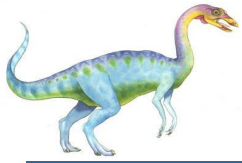
I/O-bound process
CPU-bound process





Addition of Medium Term Scheduling





Process Creation

Parent process creates **children** processes, which, in turn create other processes, forming a tree of processes.

Process identifier (pid): process identified and managed.

Resource sharing

- Parent and children share all resources

- Children share subset of parent's resources

- Parent and child share no resources

Execution

- Parent and children execute concurrently

- Parent waits until children terminate





Lecture Materials

Galvin 4.1-4.3

Galvin 13.4.1,13.4.2,13.4.3,13.4.4,13.5



End of Lecture 4, 5

