

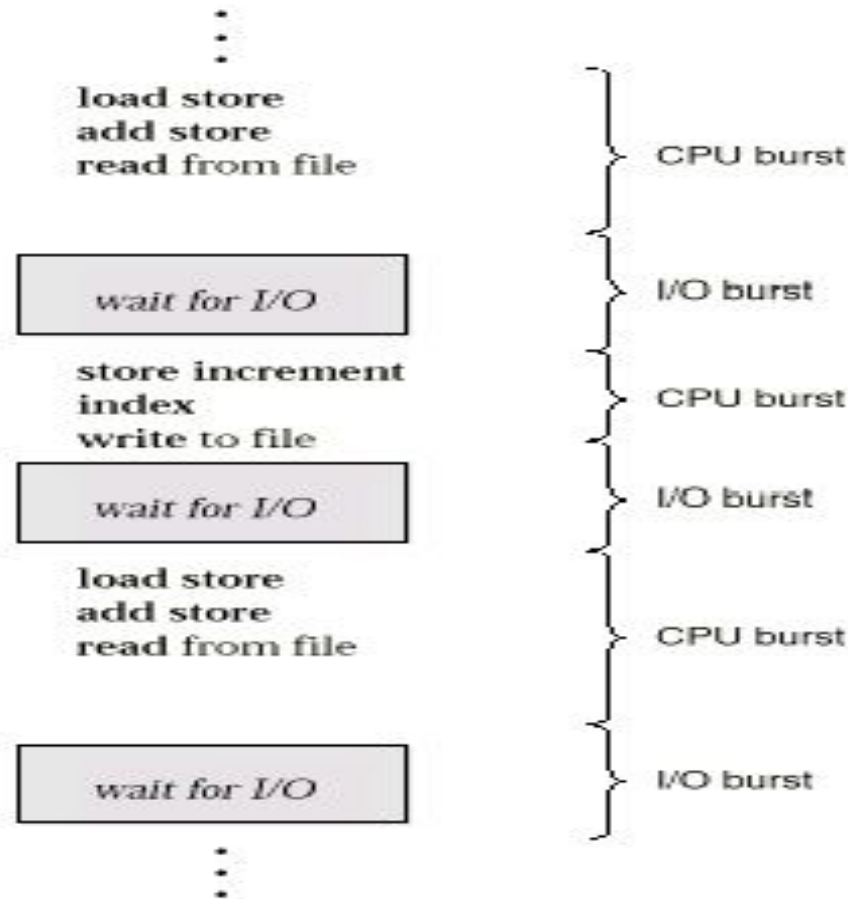
Operating Systems

CPU Scheduling Algorithms

-

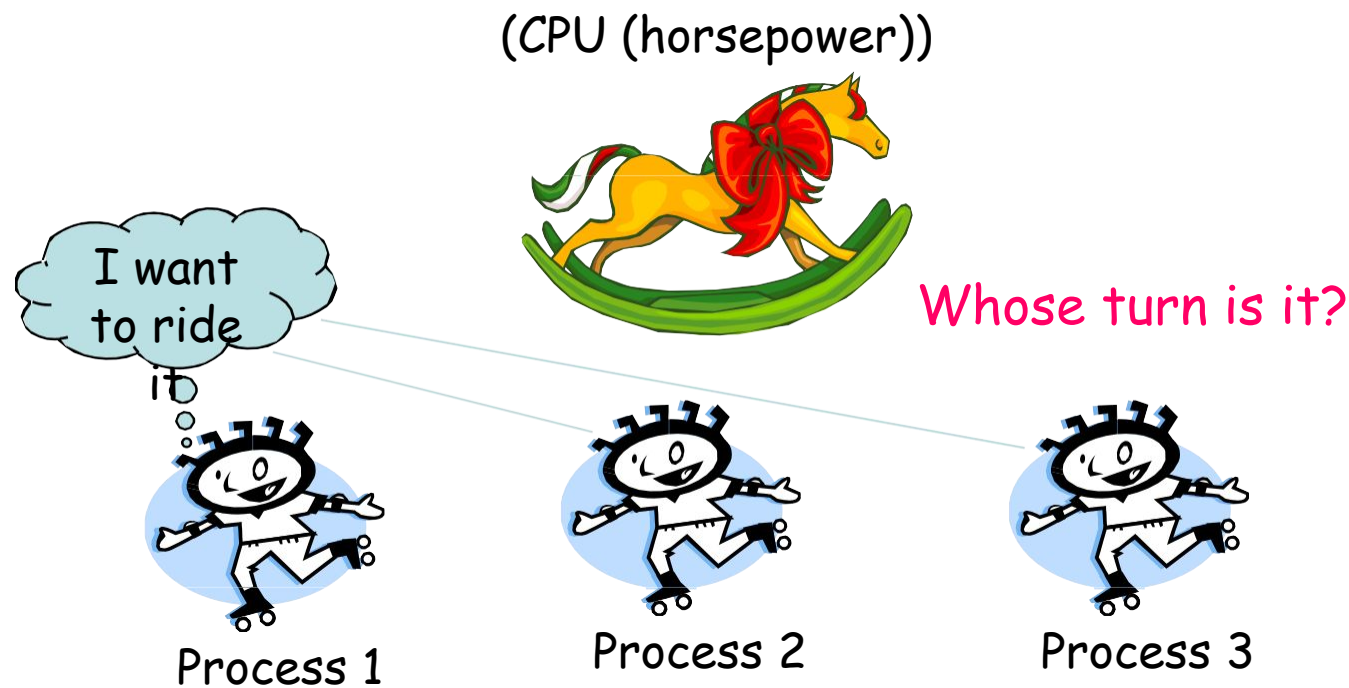
CPU – I/O Burst Cycle

- Process execution consists of a cycle of CPU execution and I/O wait
 - Processes move back & forth between these two states
- A process execution begins with a CPU burst, followed by an I/O burst and then another CPU burst and so on



Scheduling

- Deciding which process/thread should occupy a resource (CPU, disk, etc.)



CPU Scheduler

- When CPU becomes idle
 - OS must select one of the processes from the Ready Queue to be executed.
- **CPU scheduler**
 - is the OS code that implements the CPU scheduling algorithm .
 - selects a process from the processes in memory that are ready to execute (from Ready Queue), and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
 1. Switches **from running to waiting** state
 2. Switches **from running to ready** state (e.g. when time slice of a process expires or an interrupt occurs)
 3. Switches from **waiting to ready** state. (e.g. on completion of I/O)
 4. **Terminates**

Preemptive vs Non Preemptive Process

Non-preemptive

✓ Process runs until voluntarily relinquish CPU

- process blocks on an event (e.g., I/O)
- process terminates
- process periodically calls the yield() system call to give up the CPU

✓ Only suitable for domains where processes can be trusted to relinquish the CPU

Preemptive

- The scheduler actively interrupts and reschedules an executing process
- Special device requires for creating interrupt
- Required when applications cannot be trusted to yield
- Incurs some overhead

Scheduling Criteria

- **CPU utilization** – percentage of time the CPU is not idle.
- **Throughput** – completed processes per unit of time
[the amount of material or items passing through a system or process]
- **Waiting time** – time spent in the ready queue
 - For Non preemptive Algos = $\text{Starting.Time} - \text{Arrival.Time}$
 - For Preemptive Algos = $\text{Finish.Time} - \text{Arrival.Time} - \text{Burst.Time}$
- **Turnaround time** – amount of time to execute a particular process.
 - $\text{Finish.Time} - \text{Arrival.Time}$
- **Response time** – response latency
 - amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

The Perfect Scheduler

- Max CPU utilization= keep all devices busy
- Max throughput =Maximize jobs/time
- Min turnaround time
- Min waiting time
- Min response time =latency, first response
- Fairness: everyone makes progress, no one starves

Single CPU–Scheduling Algorithms

- First Come First Serve (FCFS)
- Smallest Next CPU Burst
 - Shortest Job First (SJF)
 - Shortest Remaining Time First (SRTF)
- Round Robin
- Priority Scheduling

First Come First Serve

- Simplest CPU scheduling algorithm
- Non preemptive
- The process that requests the CPU first is allocated the CPU first
- Implemented with a FIFO queue.
- **Limitations**
 - FCFS **favor long processes as compared to short ones.**
(Convoy effect)
 - **Average waiting time depends on arrival order**
 - Average waiting time is often quite long
 - FCFS is non-preemptive, so it is trouble some for time sharing systems

Convoy Effect

“A convoy effect happens when a set of processes need to use a resource for a short time, and one process holds the resource for a long time, blocking all of the other processes. Causes poor utilization of the other resources in the system”

FCFS – Example

The process that enters the ready queue first is scheduled first, regardless of the size of its next CPU burst

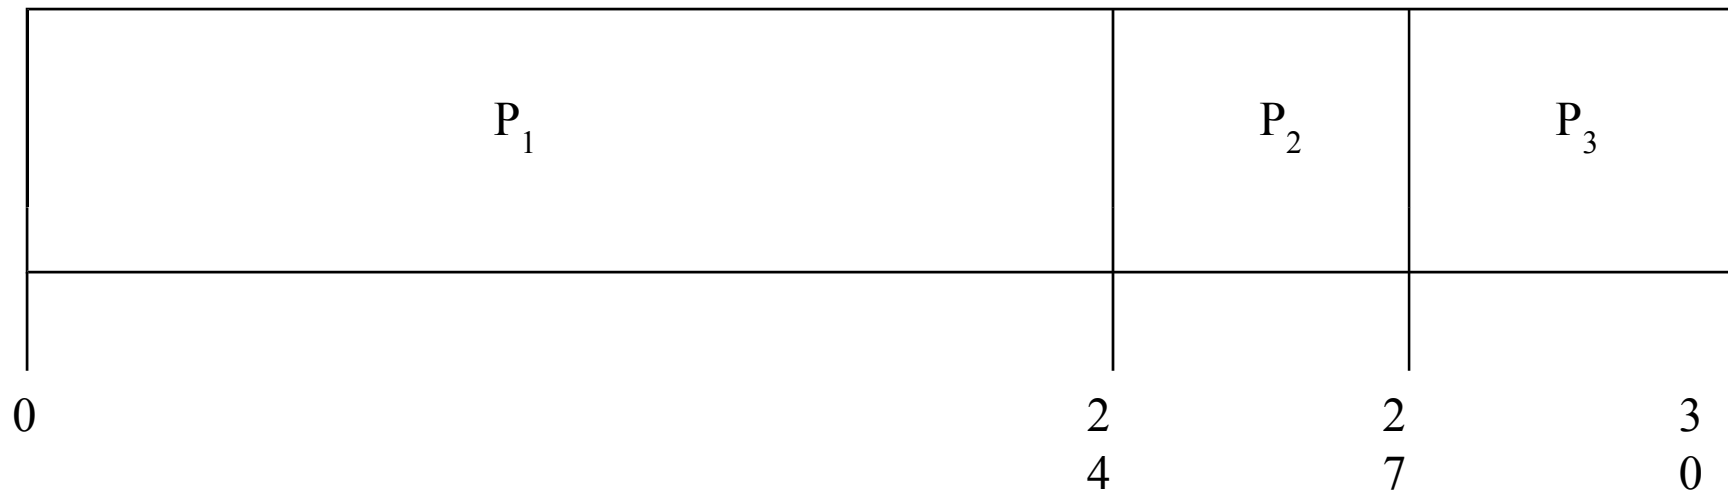
Example	Processes	Burst Time
:		24
	P	3
	2	3
	P	

Suppose that processes arrive into the system in the order: P₁, P₂, P₃

FCFS – Example

Processes are served in the order: P₁, P₂, P₃

The **Gantt Chart** for the schedule is:




Waiting times P₁ = 0; P₂ = 2; P₃ = 4

Average waiting time: $(0+2+4)/3 =$

FCFS – Example

Suppose that processes arrive in the order: P₂ , P₃ , P₁ . The Gantt chart for the schedule is:



Waiting time for P₁ = 6; P₂ = 0; P₃ = 3
Average waiting time: $(6 + 0 + 3)/3 = 3$

FCFS – Example

- Draw the graph (Gantt chart) and compute average waiting time for the following processes using **FCFS** Scheduling algorithm.

<u>Process</u>	<u>Arrival time</u>	<u>Burst Time</u>
P1	1	16
P2	5	3
P3	6	4
P4	9	2

SJF & SRTF Scheduling...

- When the CPU is available it is assigned to the process that has the smallest next CPU burst.
- If two processes have the same length next CPU bursts, FCFS scheduling is used to break the tie.

Comes in three flavors

- **Shortest Job First (SJF)**
 - It's a non preemptive algorithm.
- **Shortest Remaining Time First (SRTF)**
 - It's a Preemptive algorithm.

SJF Example

Process	Duration/B.T	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



P4 waiting time: 0-0
P1 waiting time: 3-0
P3 waiting time: 9-0
P2 waiting time: 16-0

The total running time is: 28
The average waiting time (AWT):
 $(0+3+9+16)/4 = 7$ time units

SRTF Example

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 waiting time: $12 - 0 - 10 = 2$
P2 waiting time: $4 - 2 - 2 = 0$

The average waiting time (AWT):
 $(0 + 2) / 2 = 1$

Now run this using SJF!

SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

For SJF consider all processes arrive at time 0 in sequence P1, P2, P3,

<u>P4. Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	8
P2	1	4
P3	2	9
P4	3	5

SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	5
P2	1	2
P3	2	3
P4	3	1

SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	9
P2	3	6
P3	6	2
P4	9	1

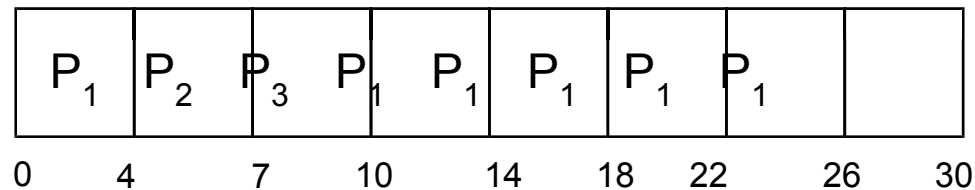
Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds.
- After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process **gets $1/n$ of the CPU time in chunks of at most q time units** at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- **Performance**
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4

<u>Proces</u>	<u>Burst</u>
<u>S</u>	<u>Time</u>
P_1	24
P_2	3
P_3	

- The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*
- q should be large compared to context switch time
- q usually 10ms to 100ms


Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority
(smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling
 - where priority is the inverse of predicted next CPU burst time
-> Shortest Job, Highest Priority
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- Priority scheduling Gantt Chart



- Non-Preemptive Waiting time = $0 + 1 + 6 + 16 + 18 = 41$
- Average waiting time = 8.2 msec

Course Materials

@Galvin-5.1-5.3